

NOT_READY_THE GLITE GATEWAY PLUGIN OF ARCLIB AND RELATED COMMAND LINE TOOLS

Technical description and user manual

Mattias Ellert¹

Ivan Marton²

Peter Stefan³

¹mattias.ellert@fysast.uu.se

²martoni@niif.hu

³stefan@niif.hu

Contents

1	Introduction	2
2	Overview	2
2.1	The supported gLite CE interface: CREAM2	2
3	External dependencies and installation instructions	3
4	ARCLIB gLite gateway plugin	3
4.1	Wrapper classes	4
4.2	Job description translators	4
4.3	CREAMClient base class	4
5	ARC command line tools for gLite CE	5
5.1	glitedelegate	5
5.2	glitesub	6
5.3	glitestat	7
5.4	glitekill	7
5.5	gliteclean	8
5.6	gliteundelegate	8
5.7	Skeleton of a temporary example gLite client	8
6	Current limitations	9
6.1	Unsupported CREAM features	9
6.2	Full ARC1 client integration	9
6.3	xRSL→JDL translation and related problems	9
6.4	RunTime Environments related issues	9
7	Conclusions	10
A	Installation from Debian binary packages	10

1 Introduction

The purpose of this document is to describe this gLite[1] specific plugin of the ARCLIB (...Hivatkozni: jvbeli Deliverable) to present the related command line utilities built on top of these libraries. The gateway library together with the tools enables seamless interoperability with gLite computing elements. It is possible, that relevant sections of this document will be moved into ARC1 userguide and into ARCLIB technical documentation. Interoperability with gLite storage element is provided by an other set of ARC library. This data library provides native interfaces to GridFTP[2] SRM which is the main building blocks of a gLite storage element. A storage interoperability is out of scope of this library.

2 Overview

Interoperability has been one of the most important disciplines in recent grid middleware development. This term is often used in the sense of accessing resources operated by one kind of grid middleware (e.g. gLite) from a user interface operated by another kind (e.g. ARC1). There are two main possibilities to achieve this result:

- either to use a special infrastructure element, the gateway, which performs a full data structure and protocol translation between the two inter-operating grid middleware solutions,
- or to explicitly implement the interface of the other side in the client.

The latter solution was preferred by ARC1 developers, the gLite gateway library module has been implemented in the ARCLIB and the gateway functionality will be transparently offered via the usual NorduGrid command line tools. An overview of the preferred solution is presented in figure 1.

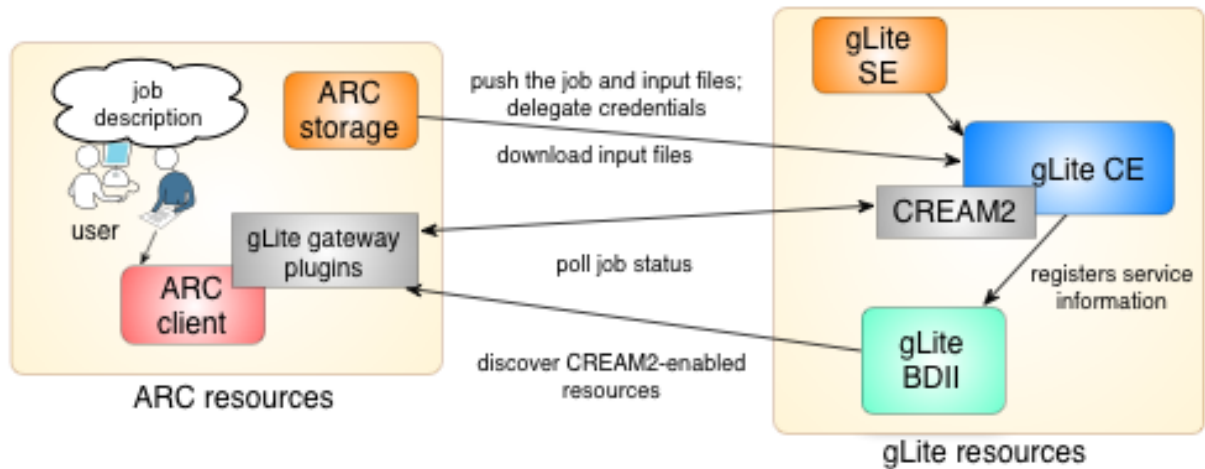


Figure 1: This chart illustrates how the ARC1→gLite gateway library module based approach works. The user will be able to express his requirements in the job description document using any of the xRSL[3], JDL or JSDL[4] languages. The ARC client library will perform an automatic transformation. The gLite gateway functionality integrated into the ARCLIB, thus into the ARC client enables the user to transparently access gLite computing resources via the usual ARC command line tools. The gLite modules perform authentication, credential delegation and job submission to the CREAM2 CE. The CREAM2 CE may perform the necessary input file download from the different storage element. The library module gets job information through CREAM2 and sends management commands, like kill a job or undelegate credentials. Note, that not all the functionality described on the overview figure is implemented yet, for the current limitations see section 6!

2.1 The supported gLite CE interface: CREAM2

The gLite computing element (CE) offers several interfaces. Some of them are only suitable for the WMS (glite resource broker) and was not meant to be used by 3rd party clients. Fortunately gLite comes with

a webservice based computing element interface, known as the CREAM CE. The CREAM has undergone several major revisions. We can only hope that some stability will soon arise. The current CREAM version has entered the EGEE certification process nevertheless it is a full scale deployment on the EGEE worldwide grid has not taken place yet. CREAM2 is available on several preproduction and testing grid sites.

Nevertheless the KnowARC project chose the CREAM version 2[5] as the gLite computing element interface. In particular, the gateway library was developed against version `glite-ce-cream_R.1.8.3.0`⁴.

CREAM2 was chosen for the following reasons:

- CREAM2 has a web service interface that fits the web service based ARC1.
- CREAM2 enables direct access to the gLite computing element without the need of going through the WMS.
- CREAM2 contains numerous improvements compared to the earlier CREAM versions.
- Unlike the other gLite CE interfaces it supports direct job status queries
- It is an official interface to gLite computing elements and as EGEE evolves it is expected that more and more computing elements will be deployed with CREAM2 interface.
- CREAM2 offers a convenient way of handling input and output files through accessing the input and output sandbox via GridFTP.

In the near future CREAM will offer yet another interface which will be based on the BES/JSDL/GLUE OGF standards. This CREAM BES interface will coexist with the CREAM2 too.

3 External dependencies and installation instructions

The ARCLIB gLite plugin have only one dependency. GridFTP components of the Globus Toolkit[6] are needed to enable data staging functionality to those CREAM2 computing elements which offer input and output sandbox over a GridFTP enabled storage. Unfortunately the gateway library only works together with a strongly patched Globus version. The NorduGrid patched Globus can be downloaded from the NorduGrid download area⁵. Further information about the NorduGrid Globus patches can be read on the NorduGrid page⁶. It is worth to mention that the user will need a VOMS[7] proxy certificate in order to use the command line tools against a CREAM2 service. This VOMS proxy can be generated by using the VOMS client tools. It is only the VOMS enabled proxy certificate that is needed here, apart from this no VOMS component needed.

The ARC gLite gateway plugin library and the temporary command line utilities are incorporated into ARC1 releases. The source code is available from the ARC1 SVN repository⁷ and binary packages are provided at the NorduGrid download area. The build and installation instructions are given in the README file distributed together with the software. A step-by-step guide for Debian binary installation can be found in the appendix A Installation from Debian binary packages.

4 ARCLIB gLite gateway plugin

The ARCLIB gLite gateway plugin consist of the wrapper classes, the job description handling facility and the gLite specific CREAMClient base class. The wrapper classes conveys the function provided by the CREAMClient to ARCLIB. The job description handling framework makes a general job request representation available for the other classes. CREAMClient base class implements the interface towards CREAM2 computing elements. This client object has an easy-to-use and intuitive interface. It uses the VOMS proxy

⁴<http://grid.pd.infn.it/cream/field.php?n=Main.ReleaseNotes>

⁵<http://download.nordugrid.org/software/globus/>

⁶http://wiki.nordugrid.org/index.php/Globus_Libraries

⁷<http://svn.nordugrid.org/repos/nordugrid/arc1/>

certificate to build the secure channel to the server and to sign that of the server's. The different functions throw *CREAMClientError* exceptions that should be caught in the application.

The following sections present the set of functions provided by the library.

4.1 Wrapper classes

... mik ezek a wrapper osztályok ... milyen base class, mit csinál, mi a funkciójuk ... viszonylag részletesen ...

4.2 Job description translators

... Brhol megváhat ... objektumot példnyostani ... milyen transzformációk llaak rendelkezésre ...

4.3 CREAMClient base class

The client object uses the VOMS proxy certificate to build the secure channel to the server and to sign that of the server's.

... 1 bekezdés arról, hogy SOAP a héd(Hivatkozni!!!) SOAP framework-jt használja 'fully integrated'... az elz deliverable-ben a kommunikáci gSOAP-ot használt, ez az egyik nagy elreps, hogy ezt lecsérltk...[8]...

CREAMClient(Arc::URL, Arc::MCCConfig)

The constructor of the class performs the necessary initialization work. It receives two arguments: the service URL and the ARC message chain component configuration file to establish the communication channel to the server. These pieces of information are indispensable to create your client and to communicate with the remote site.

setDelegationId(std::string)

Besides the constructor this is the other generic-purpose function that can be used in CREAM2 client applications. This function sets the previously registered and referred delegation ID on the client. The function is very simple: just set a private variable and returns with no value.

createDelegation(std::string)

The **createDelegation** function is about to perform the whole delegation registration process. It sends a *getProxyReq* message having the requested delegation ID, signs the received certificate and sends it back in a *putProxy* SOAP message to the server. These three steps constitute the delegation registration process. As almost every function, this one has no return value either. If there are any problems during the communication, either locally in the channel or at the remote site, the function throws *CREAMClientError* described previously. If there are no exceptions thrown, then the command is considered to have successfully completed.

destroyDelegation(std::string)

This method has the reverse functionality of the **createDelegation** function. It sends the *destroy* message of CREAM2 to the remote server.

submit(std::string)

The **submit** function is the most complex part of the class. It translates the job description received as the argument, registers the job with a *JobRegisterRequest* message, uploads the locally stored files, if necessary, then enables the job execution on the server by sending a *JobStartRequest* message. Finally it returns with a **creamJobInfo** object that contains a *jobId* (job identifier), a *creamURL* (as the service URL), an *ISB*

(reference to the Input Sandbox) and an *OSB* (Output Sandbox) member to describe the registered job. In the **glitesub** command (See Section 5.2 for details) these pieces of information are stored in the information file.

stat(std::string)

This method queries the job status from the gLite CREAM2 server. The return value is the job status translated to the ARC terminology. The possible values are presented in the section of the **glitestat** command line tool. (See Section 5.3!)

cancel(std::string)

The **cancel** function sends a *JobCancelRequest* SOAP message to the server and handles the emerging exceptions. It can be used for canceling a remotely registered and possibly running job.

purge(std::string)

The **purge** function sends a *JobPurgeRequest* message to server and throws an exception in case of any emerging problems.

5 ARC command line tools for gLite CE

The well known ARC commands (ngsub, ngget, See: The NorduGrid/ARC User Guide[9]!) will provide seamless access and integration with gLite computing elements. Full transparency will be achieved through the gLite gateway ARCLIB plugin. An ordinary user will use the same set of commands regardless of what type of CE to access. The user will be able to specify grid jobs described both in the gLite (JDL) and the NorduGrid (JSDL, xRSL) job description languages. The client will perform automatic conversion between the different job description languages.

Unfortunately the interoperability features offered by the gLite gateway library are not yet integrated into the NorduGrid commands. Also the current version of ARCLIB does not offer the full spectrum of language transformations. However only a set of temporary command line tools exposing the available features of library are provided. These temporary commands make the already existing interoperability features usable for ordinary users. There are seven different command line tools for using the gLite gateway functions. In this section these tools will be presented in detail through examples. The job descriptions and services shown are just simple examples, they might be different in real-life usage.

Every tool has a short manual page being accessible by using the **man** on-line manual reader command or some more help can be achieved by using the `<command> -?` command line option.

5.1 glitedelegate

In order to use gLite resources user must have a VOMS proxy. This proxy can be generated by command *voms-proxy-init*.

For example:

```
$ voms-proxy-init -voms knowarc.eu
```

Once a VOMS proxy has been generated everything is in place to submit a new job. If one wants submit a job to a gLite server, first has to own a valid delegation on the remote site. This can either be an old, but still valid delegation identifier, or can be a new one. There is a command line tool to register a delegation on the CE to be used.

The usage of command **glitedelegate** is straightforward. All needed is a working CREAM2 delegation service URL, and a locally unique arbitrary delegation ID. This delegation ID will be associated to the remote resource. If you are not sure whether your favourite delegation ID is occupied or not, then try to

delete it before registering it again. (See Section 5.6 for further details!) The delegation command has the following syntax:

```
$ glitedelegate <delegation ID> <delegation service URL>
```

For example:

```
$ glitedelegate test_delegation \  
    https://cream.grid.upjs.sk:8443/ce-cream/services/gridsite-delegation
```

5.2 glitesub

The current command only supports direct job submission that is the user must know service endpoints. These kind of endpoints usually stored in LDAP databases. Ordinary ldapsearch can be used to find such an endpoint. Below an example:

```
$ ldapsearch -x -h lxbra2305.cern.ch -p 2170 -b mds-vo-name=local,o=grid \  
    '(GlueCEUniqueID=*cream*)' | grep GlueCEInfoContactString | \  
    awk '{print $2}' | sort | uniq
```

The **glitesub** command implements the actual job submission in three phases: registration of a job, uploading the local files to the input sandbox URL received during the registration and starting the job.

The **glitesub** command requires a valid delegation. See Section 5.1 how to create one. By the way this is the only command which needs the delegation ID.

The next step is to prepare the necessary input files and the job description!

The glitesub command syntax is the following (please note that the service URL, job description, info file are mandatory and can only be given in the specified order):

```
$ glitesub -D <delegation ID> <service URL> <job description> <info file>
```

Here the delegation ID is the identification registered previously. The service URL will be different than in the delegation example above, because this is not the URL of the delegation, but the execution service. The job description must be in JDL format like in the example below. The info file stores information about the job submitted. It should either be a non-existing file otherwise it is overwritten. This file contains the job ID, the input sandbox URL, the output sandbox URL and the service URL.

Below I present a JDL description example. Note that the *VirtualOrganisation* and *QueueName* attributes should be the same as your virtual organisation name and the corresponding queue name on the server side.

```
[  
  Executable = "/bin/hostname";  
  StdInput = "std.in";  
  StdOutput = "std.out";  
  StdError = "std.err";  
  BatchSystem = "pbs";  
  VirtualOrganisation = "knowarc.eu";  
  InputSandbox = {"std.in"};  
  OutputSandbox = {"std.out", "std.err"};  
  QueueName = "knowarc.eu";  
  OutputSandboxDestURI = { "gsiftp://localhost/std.out", "gsiftp://localhost/std.err" };  
]
```

An example how to use this command:

```
$ glitesub -D test_delegation \
    https://cream.grid.upjs.sk:8443/ce-cream/services/CREAM2 description.jdl job.info
```

... mit csinl a parancs sikeres s sikertelen futs esetn...

5.3 glitestat

The **glitestat** command contacts to the CREAM2 service on the service URL and queries job status of a job specified by the job ID given in the info file.

Usage:

```
glitestat <info file>
```

For example:

```
glitestat job.info
```

The **glitesub** command implements a job status mapping between gLite and ARC job status semantics. The possible responses and their meanings are the following:

- ACCEPTING - the job submission is completed but the job is not yet scheduled
- SUBMITTING - scheduling in progress
- INLRMS:Q - the job is already at the local resource manager system and it is queued
- INLRMS:R - the job is waiting at the local resource manager system for running
- INLRMS:S - the job is running
- KILLED - the job was terminated
- FINISHED - the job has finished
- FAILED - the job had some failure
- FAILES - the job had some failure at LRMS level
- EXECUTED - the job has been finished and there is no state information available

5.4 glitekill

The **glitekill** command can be used for killing a remote job running in the underlying batch system. The only necessary argument needed is the information file made by **glitesub**. (See Section 5.2 for further details!) This command initiates stopping the job in the batch system. You can check the effect by using **glitestat**. (See Section 5.3 for its usage!)

Usage:

```
glitekill <info file>
```

For example:

```
glitekill job.info
```


5.5 gliteclean

Either after killing a job or because of some remote error occurs garbage files may remain on the server. These files can be purged by using command **gliteclean**. If you intentionally kill a job on the computing node, the job related stuff might also remain as garbage in the remote queue. This command also removes the job from the underlying batch system on the gLite computing element. It is important to note that this command removes the local information file as well. It might be useful to run this command after killing a job. (See Section 5.4!)

Usage:

```
gliteclean <info file>
```

For example:

```
gliteclean job.info
```

5.6 gliteundelegate

After the work is complete it is possible to unregister and delete the delegation entry from the remote site by using command **gliteundelegate**. It can also be used when one needs to register a delegation ID but are unsure whether it is already used or not. Put the delegation delegation service URL into the argument!

Usage:

```
$ gliteundelegate <delegation ID> <delegation service URL>
```

For example:

```
$ gliteundelegate test_delegation \  
    https://cream.grid.upjs.sk:8443/ce-cream/services/gridsite-delegation
```

5.7 Skeleton of a temporary example gLite client

Finally, here is an example code to present how easy and simple it is to write a new client using the CREAMClient class (See section 4.3!). This code sample requests the job status of a previously submitted remote job, and writes it to the standard output. The client is written in C++.

```
#include "CREAMClient.h"  
#include <iostream>  
  
int main(int argc, char* argv[]){  
    Arc::URL url( SERVICE_URL );  
    Arc::MCCConfig cfg;  
  
    Arc::Cream::CREAMClient gLiteClient(url,cfg);  
    try {  
        std::cout << "Job status: " << gLiteClient.stat( JOBID ) << std::endl;  
    } catch (CREAMClientException& cce) {  
        std::cerr << "ERROR: " << cce.what() << std::endl;  
        return 1;  
    }  
    return 0;  
}
```

It is of course mandatory to define *SERVICE_URL* and *JOBID* strings to their real values to reach the proper functionality.

6 Current limitations

Even though the gateway library code is ready for production use, there are a couple of issues to be handled in the upcoming development.

6.1 Unsupported CREAM features

6.2 Full ARC1 client integration

Full ARC1 client integration basically means that the previously presented commands will disappear and their functionality will be available in the standard ARC1 client toolkit. However there are two dependencies to be met to achieve this stage:

- On one hand ARC1 broker along with the BDII-capable target object should be finalized.
- On the other hand a conflicting Globus Toolkit OpenSSL and the standard OpenSSL issue is to be investigated.

The latter issue was revealed when the VOMS proxy certificate delegation method was implemented in the gateway library. Since in the ARC1 client-side library there is a built-in Message Chaining Component (MCC) which already uses Globus Toolkit OpenSSL functions to enable establishing secure communication. Linking glite* tools directly against the standard OpenSSL libs caused congesting SSL functions which led to the breakdown of the whole library. The solution had been not linking the standard OpenSSL libs, just the patched OpenSSL in the Globus Toolkit.

6.3 xRSL→JDL translation and related problems

It is possible to express jobs in xRSL which have no equivalent in JDL. Specically, in xRSL it is possible for input files downloaded from storage servers to have their names changed. This is not possible in JDL. So a user might in xRSL ask for this:

```
(inputfiles=(gsiftp://interop.dcgk.dk/storage/datafile1.txtinput.txt))
```

That is not possible to express in JDL. It only allows the file to be downloaded as datafile1.txt. This could be solved by a wrapper script that renames files according to the xRSL specification. This does not solve our problem completely however. For example if a user wants to concatenate the stdout from two jobs, he would write something like the following in xRSL:

```
(inputfiles=(gsiftp://interop.dcgk.dk/storage/job1/output.txtinput1.txt)
             (gsiftp://interop.dcgk.dk/storage/job2/output.txtinput2.txt))
```

This cannot be solved just using the renaming method, because both input files would be retrieved before renaming is done and they would therefore overwrite each other.

The solution is to recognize jobs of this type and simply barring them from using gLite resources or to let such jobs fail with an error message explaining that the user has to express their job in a different way.

6.4 RunTime Environments related issues

Many grid jobs rely on software being available at the execution location. The requirement as well as the initialization of such software is expressed through the use of Run-Time Environments (RTEs). While the library is capable of translating the RTE request itself it is not capable of translating the name of the RTE. No standardized cross Grid nomenclature has been adopted, nor are the authors aware of any such activity. A temporary solution would be for each VO to have a remotely accessible list of RTEs and their names in each job description language. This could then be used for translating jobs.

7 Conclusions

The document has described an ARC1 library extension that allows ARC1 grid users to access gLite CREAM2 computing resources. The current version of the code allows direct access to resources but as the new ARC1 brokering solution comes out the gateway functionality will be integrated fully into the ARC1 client code. The same version will also be able to access the gLite information system to extract CREAM2 resource information.

A built-in xRSL→JDL translator will also enable to use purely xRSL job description on ARC1 user interfaces.

A Installation from Debian binary packages

This section contains instructions how to install the gateway libraries and glite* command line tools on Debian Linux.

1. Put the Debian repository into the APT source files (/etc/apt/sources.list). This repository in Luebeck is maintained as long as the KnowARC project continues.

```
$ echo "deb http://pc02.inb.uni-luebeck.de:8080/~moeller/debian/stable ./" \
>> /etc/apt/sources.list
```

2. Update the package list and install the necessary dependencies:

```
$ aptitude update
$ aptitude install voms globus globus-dev globus-doc
```

3. Add the Globus libraries to the system path:

```
$ echo "/opt/globus/lib" > /etc/ld.so.conf.d/globus.conf
$ ldconfig
```

4. Install ARC1 client.

```
$ aptitude install nordugrid-arc1-client
```

5. Create the vomses file:

```
$ mkdir -p /opt/glite/etc/
```

and copy these lines into the vomses file located at /opt/glite/etc/vomses

```
"gin.ggf.org" "kuiken.nikhef.nl" "15050" "/O=dutchgrid/O=hosts/OU=nikhef.nl/CN=kuiken.nikhef.nl" "gin.ggf.org"
"knowarc.eu" "arthur.hep.lu.se" "15001" "/O=Grid/O=NorduGrid/CN=host/arthur.hep.lu.se" "knowarc.eu"
"atlas" "voms.cern.ch" "15001" "/DC=ch/DC=cern/OU=computers/CN=voms.cern.ch" "atlas"
"nordugrid.org" "voms.uninett.no" "15015" "/O=Grid/O=NorduGrid/CN=host/voms.ndgf.org" "nordugrid.org"
```

6. Create \$HOME/.globus directory and move the user certificate-key pair into there.

References

- [1] “gLite, Lightweight Middleware for Grid Computing.” [Online]. Available: <http://glite.web.cern.ch/glite/>
- [2] W. Allcock *et al.*, “Data management and transfer in high-performance computational grid environments,” *Parallel Comput.*, vol. 28, no. 5, pp. 749–771, 2002.

- [3] O. Smirnova, *Extended Resource Specification Language*, The NorduGrid Collaboration, NORDUGRID-MANUAL-4.
- [4] A. Anjomshoaa *et al.*, “Job submission description language (jsdl) specification v1.0,” December 2005, GFD-R-P.056. [Online]. Available: http://www.ggf.org/ggf_docs.final.htm
- [5] P. A. et al, “Job Submission and Management Through Web Services: the Experience with the CREAM Service.” [Online]. Available: http://grid.pd.infn.it/cream/Misc/chep07_paper.pdf
- [6] I. Foster and C. Kesselman, “Globus: A Metacomputing Infrastructure Toolkit,” *International Journal of Supercomputer Applications*, vol. 11, no. 2, pp. 115–128, 1997, available at: <http://www.globus.org>.
- [7] R. Alfieri *et al.*, “From gridmap-file to VOMS: managing authorization in a Grid environment,” *Future Gener. Comput. Syst.*, vol. 21, no. 4, pp. 549–558, 2005.
- [8] R.A. van Engelen and others, “gSOAP.” [Online]. Available: <http://www.cs.fsu.edu/~engelen/soap.html>
- [9] *The NorduGrid/ARC User Guide*, The NorduGrid Collaboration, NORDUGRID-MANUAL-6.