



NORDUGRID-XXXXXXX-NN

23/4/2009

ADMINISTRATOR'S MANUAL OF THE ARC STORAGE SYSTEM

Zsombor Nagy*

Jon Nilsen†

Salman Zubair Toor ‡

*zsombor@niif.hu

†j.k.nilsen@usit.uio.no

‡salman.toor@it.uu.se

Contents

1	Quick start guide	5
1.1	Installing a centralized system	5
1.2	Using the prototype client tool	10
1.3	Adding more storage elements	15
1.4	Make the A-Hash service replicated	22
1.5	Additional Bartenders and Librarians	25
1.6	Accessing external storage solutions	27
1.7	Using the ARC DMC	29
1.8	Using the FUSE module	30

Chapter 1

Quick start guide

1.1 Installing a centralized system

In this section we will follow the installation of the ARC storage system on a freshly installed debian lenny machine.

Get the latest ARC code from the subversion (<http://svn.nordugrid.org/>):

```
$ svn co http://svn.nordugrid.org/repos/nordugrid/arc1/trunk arc1
```

Let's check the README file for all the dependencies, and install them! We use 'aptitude' here, but of course on other linux distribution you may use a different package management software, and the name of the packages may differ as well.

```
$ sudo aptitude install build-essential
$ sudo aptitude install autoconf
$ sudo aptitude install gettext
$ sudo aptitude install cvs
$ sudo aptitude install libtool
$ sudo aptitude install pkg-config
$ sudo aptitude install libglibmm-2.4-dev
$ sudo aptitude install python-dev
$ sudo aptitude install swig
$ sudo aptitude install libxml2-dev
$ sudo aptitude install libssl-dev
```

Run the autogen and configure scripts (you may want to specify target directories for the installation with the --prefix option):

```
$ cd arc1
$ ./autogen.sh
$ ./configure --disable-java --disable-a-rex-service --disable-isi-service \
  --disable-charon-service --disable-compiler-service \
  --disable-paul-service --disable-sched-service
```

```
Unit testing:      yes
Java binding:      no
Python binding:    yes (2.5)
```

Available third-party features:

```
RLS:               no
```

```

GridFTP:          no
LFC:              no
RSL:              no
SAML:             no
MYSQL CLIENT LIB: no
gSOAP:            no

```

```

Included components:
A-Rex service:    no
ISI service:      no
CHARON service:   no
HOPI service:     yes
SCHED service:    no
STORAGE service:  yes
PAUL service:     no
SRM client (DMC): no
GSI channel (MCC): no

```

We can start to compile it:

```
$ make
```

And then install it (you need to be root if you want to install it to the default location):

```
$ sudo make install
```

Maybe we should run `ldconfig` to refresh the list of libraries:

```
$ ldconfig
```

We can check if the python packages are installed correctly:

```

$ python
Python 2.5.2 (r252:60911, Jan  4 2009, 17:40:26)
[GCC 4.3.2] on linux2
Type "help", "copyright", "credits" or "license" for more information.
>>> import arc
>>> import storage
>>> import arcom
>>>

```

If there is no error on importing any of the packages, we are good to go. If the packages cannot be imported, we should set the `PYTHONPATH` environment variable to point to their location.

It is important to synchronize the time of the machine - the storage system will make decisions based on differences between timestamps, when we deploy it on multiple machines, different times would cause problems.

We need host (and user) certificates to run the system. For testing purposes we can use the NorduGrid InstantCA solution which is capable of creating a demo-CA with short lifetime. The URL of the InstantCA service is <https://vls.grid.upjs.sk/CA/instantCA>, let's create 3 user and 3 host certificates. (e.g. we can set the Organization Name to 'knowarc', the Common name of the CA to 'storage-test', the name of the users could be 'penny', 'billy' and 'hammer', the name of the hosts could be 'upsilon', 'omicron' and 'theta', and we need a password for the CA and passwords for the user certificates, both should be at least 4 characters long).

Let's download the generated certificates and untar the archive file to see its contents:

```
$ ls -lR
.:
total 28
drwxr-xr-x 2 zsombor zsombor 4096 2009-03-19 12:11 CA
-rw-r--r-- 1 zsombor zsombor  963 2009-03-19 12:11 ca.key
-rw-r--r-- 1 zsombor zsombor  786 2009-03-19 12:11 ca.pem
drwxr-xr-x 2 zsombor zsombor 4096 2009-03-19 12:11 hostCerts
-rw-r--r-- 1 zsombor zsombor 1044 2009-03-19 12:11 readme
-rw-r--r-- 1 zsombor zsombor    3 2009-03-19 12:11 serial.srl
drwxr-xr-x 2 zsombor zsombor 4096 2009-03-19 12:11 userCerts

./CA:
total 8
-rw-r--r-- 1 zsombor zsombor 786 2009-03-19 12:11 66fe707c.0
-rw-r--r-- 1 zsombor zsombor 139 2009-03-19 12:11 66fe707c.signing_policy

./hostCerts:
total 24
-rw-r--r-- 1 zsombor zsombor 786 2009-03-19 12:11 hostcert-omicron.pem
-rw-r--r-- 1 zsombor zsombor 782 2009-03-19 12:11 hostcert-theta.pem
-rw-r--r-- 1 zsombor zsombor 786 2009-03-19 12:11 hostcert-epsilon.pem
-rw-r--r-- 1 zsombor zsombor 887 2009-03-19 12:11 hostkey-omicron.pem
-rw-r--r-- 1 zsombor zsombor 891 2009-03-19 12:11 hostkey-theta.pem
-rw-r--r-- 1 zsombor zsombor 887 2009-03-19 12:11 hostkey-epsilon.pem

./userCerts:
total 24
-rw-r--r-- 1 zsombor zsombor 778 2009-03-19 12:11 usercert-billy.pem
-rw-r--r-- 1 zsombor zsombor 778 2009-03-19 12:11 usercert-hammer.pem
-rw-r--r-- 1 zsombor zsombor 778 2009-03-19 12:11 usercert-penny.pem
-rw-r--r-- 1 zsombor zsombor 963 2009-03-19 12:11 userkey-billy.pem
-rw-r--r-- 1 zsombor zsombor 963 2009-03-19 12:11 userkey-hammer.pem
-rw-r--r-- 1 zsombor zsombor 951 2009-03-19 12:11 userkey-penny.pem
```

We have certificate and key files for all the hosts and users, and we have the CA file with the proper hashed name. There is no rule where to put these certificate, but now we will put them in `/etc/grid-security`. Let's create this directory, and a subdirectory called `certificates`. We will use the 'theta' host certificates for this machine, let's copy it there:

```
$ sudo cp hostCerts/hostcert-theta.pem /etc/grid-security/hostcert.pem
$ sudo cp hostCerts/hostkey-theta.pem /etc/grid-security/hostkey.pem
$ sudo cp CA/* /etc/grid-security/certificates/
$ ls -lR /etc/grid-security/
/etc/grid-security/:
total 12
drwxr-xr-x 2 root root 4096 2009-03-19 12:36 certificates
-rw-r--r-- 1 root root  782 2009-03-19 12:35 hostcert.pem
-rw-r--r-- 1 root root  891 2009-03-19 12:36 hostkey.pem

/etc/grid-security/certificates:
total 8
-rw-r--r-- 1 root root 786 2009-03-19 12:36 66fe707c.0
-rw-r--r-- 1 root root 139 2009-03-19 12:36 66fe707c.signing_policy
```

And I choose a user certificate and put it into the `~/.arc` directory, conveniently removing the password from the key file:

```
$ mkdir ~/.arc
```

```
$ cp userCerts/usercert-billy.pem ~/.arc/usercert.pem
$ openssl rsa -in userCerts/userkey-billy.pem -out ~/.arc/userkey.pem
$ chmod 600 ~/.arc/userkey.pem
$ ls -l ~/.arc
total 8
-rw-r--r-- 1 zsombor zsombor 778 2009-03-19 13:06 usercert.pem
-rw----- 1 zsombor zsombor 891 2009-03-19 13:09 userkey.pem
```

The storage system runs within the **arched** hosting environment daemon, which needs a configuration file which describes which services we want to run, and on which ports do we want to listen, etc. Let's copy the template configuration files to the `/etc/arc` directory:

```
$ sudo mkdir /etc/arc
$ sudo cp arc1/src/services/storage/storage_service.xml.example /etc/arc/storage_service.xml
$ sudo cp arc1/src/services/storage/clientsslconfig.xml /etc/arc
$ ls -l /etc/arc
total 12
-rw-r--r-- 1 root root 259 2009-03-19 13:26 clientsslconfig.xml
-rw-r--r-- 1 root root 4606 2009-03-19 13:25 storage_service.xml
```

The template configuration file specifies several directories where the services can store their data. We will create these directories at `/var/spool/arc`, but they can be created anywhere else.

```
$ sudo mkdir /var/spool/arc/ahash_data
$ sudo mkdir /var/spool/arc/shepherd_data
$ sudo mkdir /var/spool/arc/shepherd_store
$ sudo mkdir /var/spool/arc/shepherd_transfer
$ ls -l /var/spool/arc
total 16
drwxr-xr-x 2 root root 4096 2009-03-19 13:42 ahash_data
drwxr-xr-x 2 root root 4096 2009-03-19 13:42 shepherd_data
drwxr-xr-x 2 root root 4096 2009-03-19 13:42 shepherd_store
drwxr-xr-x 2 root root 4096 2009-03-19 13:42 shepherd_transfer
```

In this deployment we plan to run the daemon as root, so it is OK that the root owns these directories. If we create these directories somewhere else, we would need to modify the paths in the `storage_service.xml` file.

If we put the host certificates and the CA certificate somewhere else than `/etc/grid-security` then we would need to modify the `storage_service.xml` (look for `KeyPath`, `CertificatePath` and `CACertificatesDir`), and we would need to modify `/etc/arc/clientsslconfig.xml` as well. This file contains the credentials which is used by the services when they connect to an other service. If we put `clientsslconfig.xml` somewhere else than `/etc/arc` then we would need to modify all the `ClientSSLConfig` tags in the `storage_service.xml`.

Let's run the **arched** daemon with this config (specified by the `-c` option) first in the foreground (using the `-f` option) to see immediately if everything is right:

```
$ sudo /usr/local/sbin/arched -c /etc/arc/storage_service.xml -f
```

The template configuration file specifies a loglevel which only prints error messages, so nothing should be written on the screen now. The **arched** daemon should now listen on the ports: 60001 for our data transfer service called Hopi, and 60000 for all the other services:

```
$ netstat -at | grep LISTEN
tcp        0      0  *:60000          *:*               LISTEN
tcp        0      0  *:60001          *:*               LISTEN
```


Now it is time to set up or prototype client tool to access the system. It is currently called `arc_storage_cli` and it is by default installed to `/usr/local/bin`. We have to tell this CLI tool where it can find our user credentials, and a URL of a Bartender service. The Bartender service is the front-end of the storage system. We can specify these things by environment variables, or by a configuration file called `~/.arc/client.xml`. Let's create this `client.xml` file:

```
$ cat ~/.arc/client.xml
<ArcConfig>
  <KeyPath>/home/<username>/.arc/userkey.pem</KeyPath>
  <CertificatePath>/home/<username>/.arc/usercert.pem</CertificatePath>
  <CACertificatesDir>/etc/grid-security/certificates</CACertificatesDir>
  <BartenderURL>https://localhost:60000/Bartender</BartenderURL>
</ArcConfig>
```

Now we can use the `arc_storage_cli` to access our local deployment. Let's do a 'list' on the root collection ('/'):

```
$ arc_storage_cli
Usage:
  arc_storage_cli <method> [<arguments>]
Supported methods:
  stat - get detailed information about an entry
  makeCollection | make | mkdir - create a collection
  unmakeCollection | unmake | rmdir - remove an empty collection
  list | ls - list the content of a collection
  move | mv - move entries within the namespace
  putFile | put - upload a file
  getFile | get - download a file
  delFile | del | rm - remove a file
  modify | mod - modify metadata
  policy | pol - modify access policy rules
  unlink - remove a link to an entry from a collection without removing the entry itself
  credentialsDelegation | cre - delegate credentials for using gateway
  removeCredentials | rem - remove previously delegated credentials
Without arguments, each method prints its own help.
```

```
$ arc_storage_cli list
Usage: list <LN> [<LN> ...]
```

```
$ arc_storage_cli list /
- Calling the Bartender's list method...
[2009-03-19 18:37:12] [Arc.Loader] [ERROR] [12491/152991544]
  Component tcp.client(tcp) could not be created
[2009-03-19 18:37:12] [Arc.Loader] [ERROR] [12491/152991544]
  Component tls.client(tls) could not be created
[2009-03-19 18:37:12] [Arc.Loader] [ERROR] [12491/152991544]
  Component http.client(http) could not be created
[2009-03-19 18:37:12] [Arc.Loader] [ERROR] [12491/152991544]
  Component soap.client(soap) could not be created
ERROR: Wrong status from server.
Maybe wrong URL: 'https://localhost:60000/Bartender'
```

If you get this error, where the 'Component [...] could not be created', we should reinforce the python environment about the location of our python packages:

```
$ export PYTHONPATH=/usr/local/lib/python2.5/site-packages/
```

Without arguments, the `arc_storage_cli` prints the list of available methods. Specifying only a method name prints the syntax of the given method. The `list` method requires Logical Names (LN). In the storage system, every file and collection has a Logical Name, which is a path within the namespace.

Now we have a running (centralized) storage service including a Hopi service, which is a basic HTTP server, for transferring files. We have a Bartender service which is the front-end of the system to the users, so our client always connect to the Bartender service. There are several directories configured in the config file where the services store their data. These directories can always be emptied, which resets the whole system to a clean state. If we run the storage system with clean data directories, then it is completely empty, which means that there is not even a root collection.

```
$ arc_storage_cli list /
- Calling the Bartender's list method...
- done in 0.34 seconds.
'/': not found
```

The root collection has the Logical Name of '/', and it's role to provide a starting point for all the other Logical Names. We have to create this root collection first:

```
$ arc_storage_cli makeCollection /
- Calling the Bartender's makeCollection method...
- done in 0.81 seconds.
Creating collection '/': done
```

```
$ arc_storage_cli list /
- Calling the Bartender's list method...
- done in 0.55 seconds.
'/': collection
    empty.
```

Now we do have an empty root collection.

1.2 Using the prototype client tool

Assume that we have a running storage system and we want to access it with the `arc_storage_cli` client. We have the `~/.arc/client.xml` set properly, so we can list the contents of the root collection:

```
$ arc_storage_cli list /
- Calling the Bartender's list method...
- done in 0.55 seconds.
'/': collection
    empty.
```

Let's create a small file and upload it:

```
$ cat > orange
orange
$ cat orange
orange

$ arc_storage_cli put orange /
- The size of the file is 7 bytes
- Calculating md5 checksum...
- done in 0.0003 seconds, md5 checksum is 6aefd2842be62cd470709b27aedc7db7
- Calling the Bartender's putFile method...
- done in 1.26 seconds.
```

```

- Got transfer URL: http://localhost:60001/hopi/7fcef91e-ac9c-658d-4def-17d52ba33cfb
- Uploading from 'orange'
  to 'http://localhost:60001/hopi/7fcef91e-ac9c-658d-4def-17d52ba33cfb' with http...
    0 s:      0.0 kB      0.0 kB/s      0.0 kB/s      . . .
- done in 0.0182 seconds.
'orange' (7 bytes) uploaded as '/orange'.

$ arc_storage_cli list /
- Calling the Bartender's list method...
- done in 0.51 seconds.
'/': collection
  orange          <file>

```

We can get additional information about the file with the stat method:

```

$ arc_storage_cli stat /orange
- Calling the Bartender's stat method...
- done in 0.40 seconds.
'/orange': found
  states
    checksumType: md5
    neededReplicas: 1
    size: 7
    checksum: 6aefd2842be62cd470709b27aedc7db7
  entry
    owner: /DC=eu/DC=KnowARC/O=knowarc/CN=billy
    GUID: 12ce831b-eab8-1f47-29f6-19fb85ebf46a
    type: file
  parents
    0/orange: parent
  locations
    https://localhost:60000/Shepherd 3e9ea612-87c0-1e81-58f3-cec6b0b125c1: alive
  timestamps
    created: 1237504821.91

```

Here we can see the owner of the file (Billy), the location of the replicas (currently only one), the number of needed replicas (which is also: one), the checksum and size of the file, and the timestamp of the creation.

We can download this file:

```

$ arc_storage_cli get /orange /tmp
- Calling the Bartender's getFile method...
- done in 0.51 seconds.
- Got transfer URL: http://localhost:60001/hopi/521c5869-4c93-8e28-aa3a-03f477f34f63
- Downloading from 'http://localhost:60001/hopi/521c5869-4c93-8e28-aa3a-03f477f34f63'
  to '/tmp/orange' with http...
    0 s:      0.0 kB      0.0 kB/s      0.0 kB/s      . . .
- done in 1.0071 seconds.
'/orange' (7 bytes) downloaded as '/tmp/orange'.
$ cat /tmp/orange
orange

```

We can create collections which can contain files and other collections forming a tree-hierarchy. We can move files and collection around within this namespace.

```

$ arc_storage_cli makeCollection /fruits
- Calling the Bartender's makeCollection method...

```

```

- done in 0.79 seconds.
Creating collection '/fruits': done

$ arc_storage_cli makeCollection /fruits/apple
- Calling the Bartender's makeCollection method...
- done in 0.91 seconds.
Creating collection '/fruits/apple': done

$ arc_storage_cli move /orange /fruits/
- Calling the Bartender's move method...
- done in 0.99 seconds.
Moving '/orange' to '/fruits/': moved

$ arc_storage_cli list /fruits
- Calling the Bartender's list method...
- done in 0.60 seconds.
'/fruits': collection
  orange          <file>
  apple           <collection>

```

We can remove files, and remove collections (but only if the collection is empty):

```

$ arc_storage_cli rmdir /fruits
- Calling the Bartender's unmakeCollection method...
- done in 0.41 seconds.
Removing collection '/fruits': collection is not empty

$ arc_storage_cli rmdir /fruits/apple
- Calling the Bartender's unmakeCollection method...
- done in 0.89 seconds.
Removing collection '/fruits/apple': removed

$ arc_storage_cli rm /fruits/orange
- Calling the Bartender's delFile method...
- done in 0.89 seconds.
/fruits/orange: deleted

$ arc_storage_cli rmdir /fruits
- Calling the Bartender's unmakeCollection method...
- done in 0.79 seconds.
Removing collection '/fruits': removed

```

Let's introduce a new user into the system. If we have some other user certificate and key files, then we can tell `arc_storage_cli` to use them with modifying the `client.xml` or by specifying the `ARC_CERT_FILE` and `ARC_KEY_FILE` environment variables.

```

$ ARC_CERT_FILE=usercert-penny.pem ARC_KEY_FILE=userkey-penny.pem \
arc_storage_cli list /
- Calling the Bartender's list method...
- done in 0.31 seconds.
'/': denied

$ ARC_CERT_FILE=usercert-penny.pem ARC_KEY_FILE=userkey-penny.pem \
arc_storage_cli put /tmp/orange /pennys-orange
- The size of the file is 7 bytes
- Calculating md5 checksum...
- done in 0.0001 seconds, md5 checksum is 6aefd2842be62cd470709b27aadc7db7
- Calling the Bartender's putFile method...

```

```
- done in 0.81 seconds.
/pennys-orange: failed to add child to parent
```

Listing the root collection or uploading files into it for Penny is denied, because the root collection is owned by Billy, and nobody else has permissions to it. Let's ask Billy to allow for everyone to access the root collection:

```
$ arc_storage_cli pol / change ALL +read +addEntry
- Calling the Bartender's stat method...
- done in 0.30 seconds.
old action list for the user ALL was ''
the new action list of user ALL for LN / will be '+read +addEntry'
- Calling the Bartender's modify method...
- done in 0.50 seconds.
set.
```

The syntax of the policy method is the following:

```
$ arc_storage_cli pol
Usage: policy <LN> <changeType> <identity> <action list>
  <changeType> could be 'set', 'change' or 'clear'
    'set': sets the action list to the given user overwriting the old one
    'change': modify the current action list with adding and removing actions
    'clear': clear the action list of the given user
  <identity> could be a '<UserDN>' or a 'VOMS:<VO name>'
  <action list> is a list actions prefixed with '+' or '-'
    e.g. '+read +addEntry -delete'
  possible actions:
    read addEntry removeEntry delete
    modifyPolicy modifyStates modifyMetadata
```

So we have changed the action list for user 'ALL' to '+read +addEntry', which means that everybody can list the contents of this collection, create subcollection in it, upload a file into it, or move some existing entry into it.

Now we can upload a file with Penny:

```
$ ARC_CERT_FILE=usercert-penny.pem ARC_KEY_FILE=userkey-penny.pem \
arc_storage_cli put /tmp/orange /pennys-orange
- The size of the file is 7 bytes
- Calculating md5 checksum...
- done in 0.0001 seconds, md5 checksum is 6aefd2842be62cd470709b27aedc7db7
- Calling the Bartender's putFile method...
- done in 1.18 seconds.
- Got transfer URL: http://localhost:60001/hopi/4145891d-0e4c-b6bd-e4d6-dc80554b9d35
- Uploading from '/tmp/orange'
  to 'http://localhost:60001/hopi/4145891d-0e4c-b6bd-e4d6-dc80554b9d35' with http...
  1 s:      0.0 kB      0.0 kB/s      0.0 kB/s      . . .
- done in 1.0071 seconds.
'/tmp/orange' (7 bytes) uploaded as '/pennys-orange'.
```

This file now only can be downloaded by Penny, so let's allow for Billy to download it as well:

```
$ ARC_CERT_FILE=usercert-penny.pem ARC_KEY_FILE=userkey-penny.pem \
arc_storage_cli pol /pennys-orange change /DC=eu/DC=KnowARC/O=knowarc/CN=billy +read
- Calling the Bartender's stat method...
```

```

- done in 0.41 seconds.
old action list for the user /DC=eu/DC=KnowARC/O=knowarc/CN=billy was ''
the new action list of user /DC=eu/DC=KnowARC/O=knowarc/CN=billy for LN /pennys-orange
will be '+read'
- Calling the Bartender's modify method...
- done in 0.59 seconds.
set.

$ arc_storage_cli stat /pennys-orange
- Calling the Bartender's stat method...
- done in 0.43 seconds.
'/pennys-orange': found
  locations
    https://localhost:60000/Shepherd 1cc3b115-2350-8334-d206-a7b891912c00: alive
  states
    checksumType: md5
    neededReplicas: 1
    size: 7
    checksum: 6aefd2842be62cd470709b27aedc7db7
  parents
    0/pennys-orange: parent
  timestamps
    created: 1237509633.17
  policy
    /DC=eu/DC=KnowARC/O=knowarc/CN=billy: +read
  entry
    owner: /DC=eu/DC=KnowARC/O=knowarc/CN=penny
    GUID: 0e6a4c9a-d4e7-e0dc-e083-16b83d460d82
    type: file

```

We can see among the metadata of the file, that Billy has 'read' rights. Billy can now download this file:

```

$ arc_storage_cli get /pennys-orange /tmp
- Calling the Bartender's getFile method...
- done in 0.50 seconds.
- Got transfer URL: http://localhost:60001/hopi/c16d7e12-42b0-7882-c09e-1cac31dbdbcd
- Downloading from 'http://localhost:60001/hopi/c16d7e12-42b0-7882-c09e-1cac31dbdbcd'
  to '/tmp/pennys-orange' with http...
  0 s:      0.0 kB      0.0 kB/s      0.0 kB/s      . . .
- done in 1.0065 seconds.
'/pennys-orange' (7 bytes) downloaded as '/tmp/pennys-orange'.
$ cat /tmp/pennys-orange
orange

```

We can specify the needed number of replicas for a new file with the `ARC_NEEDED_REPLICAS` environment variable.

```

$ ARC_NEEDED_REPLICAS=2 arc_storage_cli put orange /new-orange
- The size of the file is 7 bytes
- Calculating md5 checksum...
- done in 0.0001 seconds, md5 checksum is 6aefd2842be62cd470709b27aedc7db7
- Calling the Bartender's putFile method...
- done in 1.11 seconds.
- Got transfer URL: http://localhost:60001/hopi/0ad54182-7618-a2af-c490-dea7c4d6fc31
- Uploading from 'orange' to
  'http://localhost:60001/hopi/0ad54182-7618-a2af-c490-dea7c4d6fc31' with http...
  1 s:      0.0 kB      0.0 kB/s      0.0 kB/s      . . .
- done in 1.0057 seconds.

```

```
'orange' (7 bytes) uploaded as '/new-orange'.

$ arc_storage_cli stat /new-orange
- Calling the Bartender's stat method...
- done in 0.40 seconds.
'/new-orange': found
  states
    checksumType: md5
    neededReplicas: 2
    size: 7
    checksum: 6aefd2842be62cd470709b27aadc7db7
  entry
    owner: /DC=eu/DC=KnowARC/O=knowarc/CN=billy
    GUID: 4e673e28-4769-50af-f02e-d7a67632f87a
    type: file
  parents
    0/new-orange: parent
  locations
    https://localhost:60000/Shepherd 4aacd276-653c-3ca5-c1a2-78a876c129c9: alive
  timestamps
    created: 1237509989.43
```

We can see that the 'neededReplica' metadata says 2, but if we deployed the system in a centralized way for testing purposes, then we only have one storage element (which is a Shepherd service and a Hopi service together), so the file can only have one replica, that's why there is only one locations.

We can modify the number of needed replicas of an existing file with the 'modify' method:

```
$ arc_storage_cli modify /new-orange set states neededReplicas 1
- Calling the Bartender's modify method...
- done in 0.63 seconds.
set

$ arc_storage_cli stat /new-orange | grep neededReplicas
    neededReplicas: 1
```

This version of the `arc_storage_cli` tool currently does not support recursive methods or working with multiple files at once.

1.3 Adding more storage elements

Let's add a new storage element to the system. A new storage element should be on a different physical machine, preferably in a different building to minimize the possibility of both storage elements being offline at the same time. But for testing purposes it is of course possible to run it on the same machine using different ports. We need to deploy a Shepherd service and a storage element service, which this time will be a 'byteio' service. The 'byteio' service implements a small part of the ByteIO interface, which means that it just put the entire file data into the SOAP message, so it is only suitable for small files.

But first modify the configuration file of our original server. Currently all the transfer URLs which were generated when we upload and download files contained the hostname 'localhost'. This is not OK if we want to access it from an other machine, we should provide transfer URLs which contains the external hostname of the machine. So lets change this line in `/etc/arc/storage_service.xml`:

```
<TURLPrefix>http://localhost:60001/hopi/</TURLPrefix>
```

to the real hostname of the machine. But this is not enough. Currently the Shepherd services uses their URLs as IDs, and the other services use these IDs to connect to a Shepherd service, so we need to change

to ID of the Shepherd service to a real external URL as well. And for the sake of clarity let's change all 'localhost' entries in the config file to our real hostname. In our case the name of this machine is just 'storage' so the Shepherd section should look like this:

```
<Service name="pythonservice" id="shepherd">
  <ClassName>storage.shepherd.shepherd.ShepherdService</ClassName>
  <ServiceID>https://storage:60000/Shepherd</ServiceID>
  <CheckPeriod>20</CheckPeriod>
  <MinCheckInterval>0.1</MinCheckInterval>
  <CreatingTimeout>600</CreatingTimeout>
  <StoreClass>arcom.store.cachedpicklestore.CachedPickleStore</StoreClass>
  <StoreCfg><DataDir>/var/spool/arc/shepherd_data</DataDir></StoreCfg>
  <BackendClass>storage.shepherd.hardlinkingbackend.HopiBackend</BackendClass>
  <BackendCfg>
    <DataDir>/var/spool/arc/shepherd_store</DataDir>
    <TransferDir>/var/spool/arc/shepherd_transfer</TransferDir>
    <TURLPrefix>http://storage:60001/hopi/</TURLPrefix>
  </BackendCfg>
  <LibrarianURL>https://storage:60000/Librarian</LibrarianURL>
  <BartenderURL>https://storage:60000/Bartender</BartenderURL>
  <ClientSSLConfig FromFile="/etc/arc/clientsslconfig.xml" />
</Service>
```

Let's check first if our centralized deployment is still working, stop the `arched` daemon remove the data directories to make it a clean start, and restart it with the modified config.

```
$ sudo rm -r /var/spool/arc/*
$ sudo /usr/local/sbin/arched -c /etc/arc/storage_service.xml -f
```

Let's see if we can still use the client:

```
$ arc_storage_cli list /
- Calling the Bartender's list method...
- done in 0.31 seconds.
'/': not found

$ arc_storage_cli makeCollection /
- Calling the Bartender's makeCollection method...
- done in 0.63 seconds.
Creating collection '/': done

$ arc_storage_cli list /
- Calling the Bartender's list method...
- done in 0.50 seconds.
'/': collection
    empty.
```

Let's upload a file and ask for two replicas:

```
$ ARC_NEEDED_REPLICAS=2 arc_storage_cli put orange /
- The size of the file is 7 bytes
- Calculating md5 checksum...
- done in 0.0005 seconds, md5 checksum is 6aefd2842be62cd470709b27aedic7db7
- Calling the Bartender's putFile method...
- done in 1.10 seconds.
- Got transfer URL: http://storage:60001/hopi/139dc57d-8c12-c188-a4a6-973449553811
- Uploading from 'orange'
```



```

    to 'http://storage:60001/hopi/139dc57d-8c12-c188-a4a6-973449553811' with http...
    1 s:      0.0 kB      0.0 kB/s      0.0 kB/s      . . .
- done in 1.0090 seconds.
'orange' (7 bytes) uploaded as '/orange'.

$ arc_storage_cli stat /orange
- Calling the Bartender's stat method...
- done in 0.41 seconds.
'/orange': found
states
  checksumType: md5
  neededReplicas: 2
  size: 7
  checksum: 6aefd2842be62cd470709b27aedic7db7
entry
  owner: /DC=eu/DC=KnowARC/O=knowarc/CN=billy
  GUID: 3fdad88a-944f-79c5-4f23-c5b3d3bc981e
  type: file
parents
  0/orange: parent
locations
  https://storage:60000/Shepherd 0b46f53c-419d-ce4e-181a-ec95d16dcc95: alive
timestamps
  created: 1237519441.96

```

Currently it has only one replica, because we still don't have more than one storage elements. While `arched` keeps running, let's create a new config for the other machine, which is called `storage2`:

```

<?xml version="1.0"?>
<ArcConfig
  xmlns="http://www.nordugrid.org/schemas/ArcConfig/2007"
  xmlns:tcp="http://www.nordugrid.org/schemas/ArcMCCTCP/2007"
>
  <Server>
    <Pidfile>/var/run/arched2.pid</Pidfile>
    <Logger level="ERROR">/var/log/arched2.log</Logger>
  </Server>
  <ModuleManager>
    <Path>/usr/local/lib/arc</Path>
  </ModuleManager>
  <Plugins><Name>mcctcp</Name></Plugins>
  <Plugins><Name>mcctls</Name></Plugins>
  <Plugins><Name>mcchttp</Name></Plugins>
  <Plugins><Name>mccsoap</Name></Plugins>
  <Chain>
    <Component name="tcp.service" id="tcp">
      <next id="tls"/>
      <tcp:Listen>
        <tcp:Port>60006</tcp:Port>
        <tcp:Version>4</tcp:Version>
      </tcp:Listen>
    </Component>
    <Component name="tls.service" id="tls">
      <next id="http"/>
      <KeyPath>/etc/grid-security/hostkey.pem</KeyPath>
      <CertificatePath>/etc/grid-security/hostcert.pem</CertificatePath>
      <CACertificatesDir>/etc/grid-security/certificates</CACertificatesDir>
    </Component>

```

```

    <Component name="http.service" id="http">
      <next id="soap">POST</next>
    </Component>
    <Component name="soap.service" id="soap">
      <next id="plexer"/>
    </Component>
    <Plexer name="plexer.service" id="plexer">
      <next id="shepherd">~/Shepherd$</next>
      <next id="byteio">~/byteio</next>
    </Plexer>
    <Service name="pythonservice" id="shepherd">
      <ClassName>storage.shepherd.shepherd.ShepherdService</ClassName>
      <ServiceID>https://storage2:60006/Shepherd</ServiceID>
      <CheckPeriod>20</CheckPeriod>
      <MinCheckInterval>0.1</MinCheckInterval>
      <CreatingTimeout>600</CreatingTimeout>
      <StoreClass>arcom.store.cachedpicklestore.CachedPickleStore</StoreClass>
      <StoreCfg>
        <DataDir>/var/spool/arc/shepherd_data2</DataDir>
      </StoreCfg>
      <BackendClass>storage.shepherd.byteio.ByteIOBackend</BackendClass>
      <BackendCfg>
        <DataDir>/var/spool/arc/shepherd_store2</DataDir>
        <TransferDir>/var/spool/arc/shepherd_transfer2</TransferDir>
        <TURLPrefix>https://storage2:60006/byteio/</TURLPrefix>
      </BackendCfg>
      <LibrarianURL>https://storage:60000/Librarian</LibrarianURL>
      <BartenderURL>https://storage:60000/Bartender</BartenderURL>
      <ClientSSLConfig FromFile="/etc/arc/clientsslconfig.xml"/>
    </Service>
    <Service name="pythonservice" id="byteio">
      <ClassName>storage.shepherd.byteio.ByteIOService</ClassName>
      <TransferDir>/var/spool/arc/shepherd_transfer2</TransferDir>
      <NotifyURL>https://storage2:60006/Shepherd</NotifyURL>
      <ClientSSLConfig FromFile="/etc/arc/clientsslconfig.xml"/>
    </Service>
  </Chain>
</ArcConfig>

```

This configuration contains only two services: a Shepherd, and byteio service. The byteio service is just a very simple storage element service which moves files within SOAP messages. The Shepherd service needs to have a backend for all supported storage element services, so it has a ByteIOBackend configured. This Shepherd connects to our original server's Bartender and Librarian services. (See the technical documentation for details about the services.)

While the original (`storage`) server still runs the services, let's start this new server:

```
$ sudo /usr/local/sbin/arched -c /etc/arc/storage_service_2.xml -f
```

And after a minute or two, let's check our file:

```

$ arc_storage_cli stat /orange
- Calling the Bartender's stat method...
- done in 0.41 seconds.
'/orange': found
states
  checksumType: md5
  neededReplicas: 2

```

```

    size: 7
    checksum: 6aefd2842be62cd470709b27aedc7db7
entry
  owner: /DC=eu/DC=KnowARC/O=knowarc/CN=billy
  GUID: 24b442eb-ab1a-680e-1f79-4c20479d3ea9
  type: file
parents
  0/orange: parent
locations
  https://storage:60000/Shepherd 315d59d3-50cd-1a10-02e6-8e6a528473c3: alive
  https://storage2:60006/Shepherd c1334c09-5d3d-983b-2081-87fe60d4764a: alive
timestamps
  created: 1237519828.4

```

It indeed has two replicas now, one on each server.

Let's download it multiple times:

```

$ arc_storage_cli get /orange /tmp/orange1
- Calling the Bartender's getFile method...
- done in 0.50 seconds.
- Got transfer URL: http://storage:60001/hoپی/de0d951f-7379-1ce4-0813-48a384c11982
- Downloading from 'http://storage:60001/hoپی/de0d951f-7379-1ce4-0813-48a384c11982'
  to '/tmp/orange1' with http...
    0 s:      0.0 kB      0.0 kB/s      0.0 kB/s      . . .
- done in 1.0060 seconds.
'/orange' (7 bytes) downloaded as '/tmp/orange1'.

$ arc_storage_cli get /orange /tmp/orange2
- Calling the Bartender's getFile method...
- done in 0.51 seconds.
- Got transfer URL: https://storage2:60006/byteio/52b6bffa-4812-03a5-93a1-834e38077ac5
- Downloading from 'https://storage2:60006/byteio/52b6bffa-4812-03a5-93a1-834e38077ac5'
  to '/tmp/orange2' with byteio...
- done in 0.2822 seconds.
'/orange' (7 bytes) downloaded as '/tmp/orange2'.

```

First it was downloaded from the **storage** server via HTTP, second time it was downloaded from the **storage2** server via ByteIO.

Let's create a third storage element on a machine called **storage3** this time using a Hopi service as storage element service like on our main **storage** server:

```

zsombor@storage:~$ cat /etc/arc/storage_service_3.xml
<?xml version="1.0"?>
<ArcConfig
  xmlns="http://www.nordugrid.org/schemas/ArcConfig/2007"
  xmlns:tcp="http://www.nordugrid.org/schemas/ArcMCCTCP/2007"
>
  <Server>
    <Pidfile>/var/run/arched3.pid</Pidfile>
    <Logger level="ERROR">/var/log/arched3.log</Logger>
  </Server>
  <ModuleManager>
    <Path>/usr/local/lib/arc/</Path>
  </ModuleManager>
  <Plugins><Name>mcctcp</Name></Plugins>
  <Plugins><Name>mcctls</Name></Plugins>
  <Plugins><Name>mcchttp</Name></Plugins>

```

```

<Plugins><Name>mccsoap</Name></Plugins>
<Chain>
  <Component name="tcp.service" id="tcp">
    <next id="tls"/>
    <tcp:Listen>
<tcp:Port>60010</tcp:Port>
<tcp:Version>4</tcp:Version>
    </tcp:Listen>
  </Component>
  <Component name="tls.service" id="tls">
    <next id="http"/>
    <KeyPath>/etc/grid-security/hostkey.pem</KeyPath>
    <CertificatePath>/etc/grid-security/hostcert.pem</CertificatePath>
    <CACertificatesDir>/etc/grid-security/certificates</CACertificatesDir>
  </Component>
  <Component name="http.service" id="http">
    <next id="soap">POST</next>
  </Component>
  <Component name="soap.service" id="soap">
    <next id="plexer"/>
  </Component>
  <Plexer name="plexer.service" id="plexer">
    <next id="shepherd">~/Shepherd$</next>
  </Plexer>
  <Service name="pythonservice" id="shepherd">
    <ClassName>storage.shepherd.shepherd.ShepherdService</ClassName>
    <ServiceID>https://storage3:60010/Shepherd</ServiceID>
    <CheckPeriod>20</CheckPeriod>
    <MinCheckInterval>0.1</MinCheckInterval>
    <CreatingTimeout>600</CreatingTimeout>
    <StoreClass>arcom.store.cachedpicklestore.CachedPickleStore</StoreClass>
    <StoreCfg><DataDir>/var/spool/arc/shepherd_data3</DataDir></StoreCfg>
    <BackendClass>storage.shepherd.hardlinkingbackend.HopiBackend</BackendClass>
    <BackendCfg>
      <DataDir>/var/spool/arc/shepherd_store3</DataDir>
      <TransferDir>/var/spool/arc/shepherd_transfer3</TransferDir>
      <TURLPrefix>http://storage3:60011/hoi/</TURLPrefix>
    </BackendCfg>
    <LibrarianURL>https://storage:60000/Librarian</LibrarianURL>
    <BartenderURL>https://storage:60000/Bartender</BartenderURL>
    <ClientSSLConfig FromFile="/etc/arc/clientsslconfig.xml" />
  </Service>
</Chain>
<Chain>
  <Component name="tcp.service" id="tcp2">
    <next id="http2"/>
    <tcp:Listen>
<tcp:Port>60011</tcp:Port>
<tcp:Version>4</tcp:Version>
    </tcp:Listen>
  </Component>
  <Component name="http.service" id="http2">
    <next id="plexer2">GET</next>
    <next id="plexer2">PUT</next>
  </Component>
  <Plexer name="plexer.service" id="plexer2">
    <next id="hoi">~/hoi/</next>
  </Plexer>

```

```

    <Service name="hopi" id="hopi">
      <DocumentRoot>/var/spool/arc/shepherd_transfer3</DocumentRoot>
      <SlaveMode>1</SlaveMode>
    </Service>
  </Chain>
</ArcConfig>

```

This configuration has two separate ports, one with TLS security and a Shepherd service, and the other without TLS security and with the Hopi service. Let's run this server as well:

```
$ sudo /usr/local/sbin/arched -c /etc/arc/storage_service_3.xml -f
```

Now when checking our file again there should be no change at all:

```
$ arc_storage_cli stat /orange
locations
https://storage:60000/Shepherd 315d59d3-50cd-1a10-02e6-8e6a528473c3: alive
https://storage2:60006/Shepherd c1334c09-5d3d-983b-2081-87fe60d4764a: alive
```

Let's stop our server on `storage2`, and watch what happens of our file:

```
$ arc_storage_cli stat /orange
locations
https://storage:60000/Shepherd 315d59d3-50cd-1a10-02e6-8e6a528473c3: alive
https://storage2:60006/Shepherd c1334c09-5d3d-983b-2081-87fe60d4764a: offline

$ arc_storage_cli stat /orange
locations
https://storage:60000/Shepherd 315d59d3-50cd-1a10-02e6-8e6a528473c3: alive
https://storage2:60006/Shepherd c1334c09-5d3d-983b-2081-87fe60d4764a: offline
https://storage3:60010/Shepherd 421aa939-81fa-5cd9-7c2f-2d5f99a98be1: creating

$ arc_storage_cli stat /orange
locations
https://storage:60000/Shepherd 315d59d3-50cd-1a10-02e6-8e6a528473c3: alive
https://storage2:60006/Shepherd c1334c09-5d3d-983b-2081-87fe60d4764a: offline
https://storage3:60010/Shepherd 421aa939-81fa-5cd9-7c2f-2d5f99a98be1: alive
```

A new replica have been created on the third server to maintain the needed number of 2.

Let's start the services on `storage2` again, and let's see what is happening with our file:

```
$ arc_storage_cli stat /orange
locations
https://storage:60000/Shepherd 315d59d3-50cd-1a10-02e6-8e6a528473c3: alive
https://storage2:60006/Shepherd c1334c09-5d3d-983b-2081-87fe60d4764a: alive
https://storage3:60010/Shepherd 421aa939-81fa-5cd9-7c2f-2d5f99a98be1: alive

$ arc_storage_cli stat /orange
locations
https://storage:60000/Shepherd 315d59d3-50cd-1a10-02e6-8e6a528473c3: alive
https://storage2:60006/Shepherd c1334c09-5d3d-983b-2081-87fe60d4764a: thirdwheel
https://storage3:60010/Shepherd 421aa939-81fa-5cd9-7c2f-2d5f99a98be1: alive

$ arc_storage_cli stat /orange
locations
https://storage:60000/Shepherd 315d59d3-50cd-1a10-02e6-8e6a528473c3: alive
https://storage3:60010/Shepherd 421aa939-81fa-5cd9-7c2f-2d5f99a98be1: alive
```

One of the replicas have been disappeared to maintain the needed number of 2.

1.4 Make the A-Hash service replicated

Now we have a deployment, where one machine has the A-Hash, Librarian and Bartender services, and all three machines have a Shepherd and a storage element service. The A-Hash service stores all the metadata of all the files and collections, so it is critical. If the first machine goes offline then the whole system dies. We can deploy the A-Hash service on all the machines and still can have one consistent metadata database. The replicated version of the A-Hash uses the Berkeley DB to provide a distributed database.

Let's install Berkeley DB as root (but you can use the `--prefix` option of `configure` to install it somewhere else than the default location). This should be done on all the machines:

Get the latest Berkeley DB and its patches:

```
$ wget http://download.oracle.com/berkeley-db/db-4.7.25.tar.gz
$ wget http://www.oracle.com/technology/products/berkeley-db/db/update/4.7.25/patch.4.7.25.1
$ wget http://www.oracle.com/technology/products/berkeley-db/db/update/4.7.25/patch.4.7.25.2
$ wget http://www.oracle.com/technology/products/berkeley-db/db/update/4.7.25/patch.4.7.25.3
```

Extract it and apply the patches:

```
$ tar xzf db-4.7.25.tar.gz
$ cd db-4.7.25
$ patch -p 0 < ../patch.4.7.25.1
$ patch -p 0 < ../patch.4.7.25.2
$ patch -p 0 < ../patch.4.7.25.3
```

Configure, make and install it (you can use the `--prefix` option for `configure`)

```
$ cd build_unix/
$ ../dist/configure
```

Now compile and install it:

```
$ make
$ sudo make install
```

Now we need the python wrapper for the Berkeley DB:

```
$ wget http://pypi.python.org/packages/source/b/bsddb3/bsddb3-4.7.5.tar.gz
```

Let's install it to the default location:

```
$ tar xzf bsddb3-4.7.5.tar.gz
$ cd bsddb3-4.7.5
$ sudo python setup.py install
```

If you've installed Berkeley DB somewhere else, or want to install the python wrapper to somewhere else than the default:

```
$ BERKELEYDB_DIR=/path/to/berkeleydb python setup.py install --prefix=/where/to/install
```

Let's test the installation with the `test.py` included in the directory of the python wrapper for the Berkeley DB:

```
$ sudo python test.py
[...]
Ran 383 tests in 29.858s
```

OK

Let's configure our first server (**storage**). Here is the AHash part of the config:

```
<Service name="pythonservice" id="ahash">
  <ClassName>storage.ahash.ahash.AHashService</ClassName>
  <AHashClass>storage.ahash.replicatedahash.ReplicatedAHash</AHashClass>
  <LocalDir>/var/spool/arc/ahash_data</LocalDir>
  <MyURL>https://storage:60000/AHash</MyURL>
  <OtherURL>https://storage2:60006/AHash</OtherURL>
  <Priority>50</Priority>
  <ClientSSLConfig FromFile='/etc/arc/clientsslconfig.xml' />
</Service>
```

In this config we specify the second server as a peer. Let's configure the second server (**storage2**) as well:

```
<Plexer name="plexer.service" id="plexer">
  <next id="ahash">~/AHash$</next>
  <next id="shepherd">~/Shepherd$</next>
  <next id="byteio">~/byteio</next>
</Plexer>
<Service name="pythonservice" id="ahash">
  <ClassName>storage.ahash.ahash.AHashService</ClassName>
  <AHashClass>storage.ahash.replicatedahash.ReplicatedAHash</AHashClass>
  <LocalDir>/var/spool/arc/ahash_data2</LocalDir>
  <MyURL>https://storage2:60006/AHash</MyURL>
  <OtherURL>https://storage:60000/AHash</OtherURL>
  <Priority>50</Priority>
  <ClientSSLConfig FromFile='/etc/arc/clientsslconfig.xml' />
</Service>
```

In this config we specify the first server as a peer. So the two together will be able to form a replicated network with two nodes. But let's configure the third one (**storage3**):

```
<Plexer name="plexer.service" id="plexer">
  <next id="ahash">~/AHash$</next>
  <next id="shepherd">~/Shepherd$</next>
</Plexer>
<Service name="pythonservice" id="ahash">
  <ClassName>storage.ahash.ahash.AHashService</ClassName>
  <AHashClass>storage.ahash.replicatedahash.ReplicatedAHash</AHashClass>
  <LocalDir>/var/spool/arc/ahash_data3</LocalDir>
  <MyURL>https://storage3:60010/AHash</MyURL>
  <OtherURL>https://storage:60000/AHash</OtherURL>
  <Priority>50</Priority>
  <ClientSSLConfig FromFile='/etc/arc/clientsslconfig.xml' />
</Service>
```

Here we specify the first server as the peer.

It is important that we need to make a fresh start on all the servers, and remove the existing database, which was created by a different store:

```
$ sudo rm -r /var/spool/arc/*
```

Let's start the first server. It will complain that it cannot find its peer:

```
ERROR connecting to https://storage2:60006/AHash
[2009-03-24 10:34:18] [Arc.Storage.ReplicatedAHash] [ERROR] [10827/156026112]
    failed to send to 2 of [1, 2]
[2009-03-24 10:34:20] [Arc.Storage.ReplicatedAHash] [ERROR] [10827/158713512]
    Couldn't run election
```

With only one node of the replicated A-Hash the system is not ready to work, if we try to use it, we get an error message:

```
$ arc_storage_cli makeCollection /
- Calling the Bartender's makeCollection method...
- done in 0.55 seconds.
Creating collection '/': failed to create new librarian entry
```

Let's start the second server. It may first print some error messages, but then both servers should be quiet, and now we can use the system. Let's upload a file, and set the number of needed replicas to 3:

```
$ arc_storage_cli makeCollection /
- Calling the Bartender's makeCollection method...
- done in 1.16 seconds.
Creating collection '/': done
$ arc_storage_cli put orange /
- The size of the file is 7 bytes
- Calculating md5 checksum...
- done in 0.0214 seconds, md5 checksum is 6aefd2842be62cd470709b27aedc7db7
- Calling the Bartender's putFile method...
- done in 2.22 seconds.
- Got transfer URL: http://storage:60001/hopi/3f903c50-0424-2891-5730-521cc00d96dc
- Uploading from 'orange'
  to 'http://storage:60001/hopi/3f903c50-0424-2891-5730-521cc00d96dc' with http...
  1 s:      0.0 kB      0.0 kB/s      0.0 kB/s      . . .
- done in 1.0293 seconds.
'orange' (7 bytes) uploaded as '/orange'.
$ arc_storage_cli modify /orange set states neededReplicas 3
- Calling the Bartender's modify method...
- done in 0.74 seconds.
set
$ arc_storage_cli stat /orange
[...]
locations
  https://storage:60000/Shepherd f4ad8530-354d-5066-04ae-c54c94fd525b: alive
[...]
zsombor@storage:~$ arc_storage_cli stat /orange
[...]
locations
  https://storage:60000/Shepherd f4ad8530-354d-5066-04ae-c54c94fd525b: alive
  https://storage3:60010/Shepherd 3c55924b-73ad-c226-0f4a-4c056c967314: alive
  https://storage2:60006/Shepherd 97edebb5-e72d-098a-55a7-14595f668770: alive
[...]
```

It cannot create 3 replicas, because are third server is not running. Let's start it. After a minute or two, or file has three replicas:

```
$ arc_storage_cli stat /orange
[...]
locations
  https://storage:60000/Shepherd f4ad8530-354d-5066-04ae-c54c94fd525b: alive
  https://storage3:60010/Shepherd 3c55924b-73ad-c226-0f4a-4c056c967314: alive
  https://storage2:60006/Shepherd 97edebb5-e72d-098a-55a7-14595f668770: alive
[...]
```

The fact that this file has three replicas is not at all related to the number of AHashes, but to the number of storage elements (Shepherds and storage element services), so this is no different then what we saw in the last section. The benefit of having more then one AHashes in a replicated way is that if one of the machines goes offline, the AHash is still working. But in our current case if we shutdown the first server, while the

AHash will still work but we still have only one instance of the other two important service: the Librarian and the Bartender, so the whole system would be unusable. Let's add more Librarians and Bartenders to the system.

1.5 Additional Bartenders and Librarians

Let's add a Bartender and a Librarian to each servers. We don't need to modify our first server now. For the other two we need to change the `Plexer` section, add the two services, and modify the Shepherd to use the server's own Bartender and Librarian for replication and heartbeats. Let's modify the config of our second server: ([...] indicates unchanged parts)

```
[...]
<Plexer name="plexer.service" id="plexer">
  <next id="ahash">~/AHash$</next>
  <next id="shepherd">~/Shepherd$</next>
  <next id="librarian">~/Librarian$</next>
  <next id="bartender">~/Bartender$</next>
  <next id="byteio">~/byteio</next>
</Plexer>
<Service name="python.service" id="librarian">
  <ClassName>storage.librarian.librarian.LibrarianService</ClassName>
  <AHashURL>https://storage2:60006/AHash</AHashURL>
  <HeartbeatTimeout>30</HeartbeatTimeout>
  <CheckPeriod>20</CheckPeriod>
  <ClientSSLConfig FromFile="/etc/arc/clientsslconfig.xml" />
</Service>
<Service name="python.service" id="bartender">
  <ClassName>storage.bartender.bartender.BartenderService</ClassName>
  <LibrarianURL>https://storage2:60006/Librarian</LibrarianURL>
  <ClientSSLConfig FromFile="/etc/arc/clientsslconfig.xml" />
</Service>
[...]
<Service name="python.service" id="shepherd">
[...]
```

And the third server:

```
[...]
<Plexer name="plexer.service" id="plexer">
  <next id="ahash">~/AHash$</next>
  <next id="shepherd">~/Shepherd$</next>
  <next id="librarian">~/Librarian$</next>
  <next id="bartender">~/Bartender$</next>
</Plexer>
<Service name="python.service" id="librarian">
  <ClassName>storage.librarian.librarian.LibrarianService</ClassName>
  <AHashURL>https://storage3:60010/AHash</AHashURL>
  <HeartbeatTimeout>30</HeartbeatTimeout>
  <CheckPeriod>20</CheckPeriod>
  <ClientSSLConfig FromFile="/etc/arc/clientsslconfig.xml" />
</Service>
<Service name="python.service" id="bartender">
  <ClassName>storage.bartender.bartender.BartenderService</ClassName>
  <LibrarianURL>https://storage3:60010/Librarian</LibrarianURL>
```

```

    <ClientSSLConfig FromFile="/etc/arc/clientsslconfig.xml" />
</Service>
[...]
<Service name="pythonservice" id="shepherd">
[...]
    <LibrarianURL>https://storage3:60010/Librarian</LibrarianURL>
    <BartenderURL>https://storage3:60010/Bartender</BartenderURL>
[...]

```

We need to modify our `client.xml` to add all the other Bartenders:

```

$ cat ~/.arc/client.xml
<ArcConfig>
  <KeyPath>/home/zsombor/.arc/userkey.pem</KeyPath>
  <CertificatePath>/home/zsombor/.arc/usercert.pem</CertificatePath>
  <CACertificatesDir>/etc/grid-security/certificates</CACertificatesDir>
  <BartenderURL>https://storage:60000/Bartender</BartenderURL>
  <BartenderURL>https://storage2:60006/Bartender</BartenderURL>
  <BartenderURL>https://storage3:60010/Bartender</BartenderURL>
</ArcConfig>

```

Now we can start all three servers. The replicated A-Hash sometimes prints great amount of error messages, but if it stops printing them, then everything is OK. Now we can use the system as intended. There is a flag for the `arc_storage_cli` which makes it very verbose. We can check the used Bartender with it:

```

$ arc_storage_cli -v stat /orange
- The URL of the Bartender(s): https://storage:60000/Bartender,
  https://storage2:60006/Bartender, https://storage3:60010/Bartender
- The key file: /home/zsombor/.arc/userkey.pem
- The cert file: /home/zsombor/.arc/usercert.pem
- The CA dir: /etc/grid-security/certificates
- Calling the Bartender's stat method...
- trying https://storage:60000/Bartender
[...]
- done in 1.27 seconds.
'/orange': found

```

Now we can stop our first server. The system will be unaccessible for a few seconds while the other two A-Hash figure out what is happening, but after that it will work again:

```

$ arc_storage_cli -v stat /orange
- The URL of the Bartender(s): https://storage:60000/Bartender,
  https://storage2:60006/Bartender, https://storage3:60010/Bartender
- The key file: /home/zsombor/.arc/userkey.pem
- The cert file: /home/zsombor/.arc/usercert.pem
- The CA dir: /etc/grid-security/certificates
- Calling the Bartender's stat method...
- trying https://storage:60000/Bartender
[...]
- (-1, "Wrong status from server.\nMaybe wrong URL: 'https://storage:60000/Bartender'")
- trying https://storage2:60006/Bartender
[...]
- done in 0.45 seconds.
'/orange': found
[...]

```

If we start again our first server, after a few seconds the system will work again with the three fully-functional node.

```

zsombor@storage:~$ arc_storage_cli -v stat /orange
- The URL of the Bartender(s): https://storage:60000/Bartender,
  https://storage2:60006/Bartender, https://storage3:60010/Bartender
- The key file: /home/zsombor/.arc/userkey.pem
- The cert file: /home/zsombor/.arc/usercert.pem
- The CA dir: /etc/grid-security/certificates
- Calling the Bartender's stat method...
- trying https://storage:60000/Bartender
[...]
- done in 0.49 seconds.
'/orange': found
  states
    checksumType: md5
    neededReplicas: 3
    size: 7
    checksum: 6aefd2842be62cd470709b27aedc7db7
  entry
    owner: /DC=eu/DC=KnowARC/O=knowarc/CN=billy
    GUID: 0080d8be-1f6f-2fc7-ad71-732f899326c8
    type: file
  parents
    0/orange: parent
  locations
    https://storage3:60010/Shepherd 03312843-d3b3-2d15-4732-ad673df958ae: alive
    https://storage2:60006/Shepherd b0057765-4532-d0b4-fded-182499a91450: alive
    https://storage:60000/Shepherd 532ec0f8-dd83-297b-d7b6-dc5dacea34e0: alive
  timestamps
    created: 1237887474.95

```

1.6 Accessing external storage solutions

If we have permissions to access some third-party storage solutions (e.g. dCache via the GridFTP protocol), we can create a mount point within the namespace of the ARC storage system which points to a third-party URL, and then we can use the client interface of the ARC storage to access the third-party storage as well. This is achieved by a module of the Bartender, called the 'Gateway'. Currently only GridFTP access is supported, and for that we need to install the Globus libraries to our machines. Then we need to recompile our source for GridFTP support. The output of the `configure` script should look like this:

Available third-party features:

```

RLS:                yes
GridFTP:             yes
LFC:                 no
RSL:                 yes
SAML:                yes
MYSQL CLIENT LIB:    no
gSOAP:               yes

```

Included components:

```

A-Rex service:       no
ISI service:         no
CHARON service:      no
HOPI service:        yes
SCHED service:       no
STORAGE service:     yes
PAUL service:        no
SRM client (DMC):    yes

```

```
GSI channel (MCC): yes
```

After compilation and installation we can use the `arcls` tool to list the content of a third-party storage. Let's check if our Billy user has access to this directory:

```
$ arcls gsiftp://sal1.uppmax.uu.se:2811/pnfs/uppmax.uu.se/data
data1
data2
other
```

We can see some files in there, so we have permissions to list that directory.

Now let's configure the Bartenders on all three machines to use the Gateway module:

```
<Service name="pythonservice" id="bartender">
  <ClassName>storage.bartender.bartender.BartenderService</ClassName>
  <LibrarianURL>[...]</LibrarianURL>
  <ProxyStore>/var/spool/arc/proxy_store</ProxyStore>
  <GatewayClass>storage.bartender.gateway.gateway.Gateway</GatewayClass>
  <GatewayCfg>
    <ProxyStore>/var/spool/arc/proxy_store</ProxyStore>
    <CACertificatesDir>/etc/grid-security/certificates</CACertificatesDir>
  </GatewayCfg>
  <ClientSSLConfig FromFile="/etc/arc/clientsslconfig.xml" />
</Service>
```

Let's start all three servers. Then create a mount point with the Logical Name `/salmount` which points to this URL:

```
$ arc_storage_cli makemount /salmount \
  gsiftp://sal1.uppmax.uu.se:2811/pnfs/uppmax.uu.se/data
- Calling the Bartender's makeMountpoint method...
- done in 3.07 seconds.
Creating mountpoint '/salmount': done

$ arc_storage_cli stat /salmount
- Calling the Bartender's stat method...
- done in 0.47 seconds.
'/salmount': found
states
  closed: 0
entry
  owner: /DC=eu/DC=KnowARC/O=knowarc/CN=billy
  type: mountpoint
  GUID: b21d9d22-e885-502a-ee0e-bd0a7b967543
parents
  0/salmount: parent
mountpoint
  externalURL: gsiftp://sal1.uppmax.uu.se:2811/pnfs/uppmax.uu.se/data
timestamps
  created: 1237990168.09
```

Let's list the content of this mountpoint:

```
$ arc_storage_cli list /salmount- Calling the Bartender's list method...
- done in 0.49 seconds.
'/salmount': Your proxy cannot be found. Please delegate your credentials!
```

Yes, we first need to delegate our credentials with the `credentialsDelegation` method, then try to list the mount point:

```
$ arc_storage_cli cre
- Calling the Bartender's credentialsDelegation method...
[...]
- done in 0.42 seconds.
Successful delegation.
Proxy ID: eNCMDmi8i6YnPSAtDmVmuSEmABFKDmABFKDmOzKKDmFBFKDmeOzdBo

$ arc_storage_cli list /salmount
- Calling the Bartender's list method...
- done in 3.19 seconds.
'/salmount': collection
  data1 <unknown>
  data2 <unknown>
  other <unknown>
```

Let's download a file:

```
$ arc_storage_cli get /salmount/data1 /tmp
- Calling the Bartender's getFile method...
- done in 3.56 seconds.
- Got transfer URL: gsiftp://sal1.uppmax.uu.se:2811/pnfs/uppmax.uu.se/data/data1
- Downloading from 'gsiftp://sal1.uppmax.uu.se:2811/pnfs/uppmax.uu.se/data/data1'
  to '/tmp/data1' with gridftp...
   2 s:          0.0 kB          0.0 kB/s          0.0 kB/s      . . .
   3 s:          1.3 kB          0.4 kB/s          0.4 kB/s      . . .
- done in 2.7828 seconds.
'/salmount/data1' (1290 bytes) downloaded as '/tmp/data1'.
```

1.7 Using the ARC DMC

ARC has its own multi-protocol storage client tools which provide a unified way to access different kind of storage. For all the supported types there is a DMC module. We created a DMC module for the ARC storage as well. If it is compiled and installed, then you can use the `arcls`, `arccp`, `arcrm`, etc. tools with the `arc://` protocol to access the ARC storage system. You can specify the URL of the Bartender with the `BartenderURL` URL option, or you can have a Bartender URL configured in the `client.xml` user config.

```
$ arcls arc:///
$ arccp orange arc:///orange
$ arcls arc:///
orange
$ arcls arc:///BartenderURL=https://storage2:60006/Bartender
orange
$ arc_storage_cli list /
- Calling the Bartender's list method...
- done in 0.56 seconds.
'/': collection
  orange <file>
$ arc_storage_cli makeCollection /fruits
- Calling the Bartender's makeCollection method...
- done in 2.38 seconds.
Creating collection '/fruits': done
$ arccp arc:///orange arc:///fruits/orange
$ arcls arc:///
```

```

orange
fruits
$ arcls arc:///fruits
orange
$ arccp arc:///orange /tmp/orange
$ cat /tmp/orange
orange
$ arcrm arc:///orange
$ arcls arc:///
fruits

```

1.8 Using the FUSE module

Besides the `arc_storage_cli` tool we can access the storage system with the use of a FUSE module. We can install the ARC storage FUSE module following these steps:

- Install ARC1 storage system.
- Install `fuse >= 2.7.3` and `fuse-python >= 0.2`:

```
$ sudo aptitude install fuse-python
```

- Set up `~/.arc/client.xml`:

```

$ cat ~/.arc/client.xml
<ArcConfig>
  <KeyPath>/home/<username>/.arc/userkey.pem</KeyPath>
  <CertificatePath>/home/<username>/.arc/usercert.pem</CertificatePath>
  <CACertificatesDir>/etc/grid-security/certificates</CACertificatesDir>
  <BartenderURL>https://localhost:60000/Bartender</BartenderURL>
</ArcConfig>

```

- Create a mountpoint and mount `arcfs.py`:

```

$ pwd
/home/me/arc
$ ln -s arc1/src/services/storage/fuse/arcfs.py arcfs.py
$ mkdir mnt
$ python arcfs.py ./mnt

```

This mounts ARCFs in `/home/me/arc/mnt`, creates directory `/home/me/arc/fuse_transfer` and a log file `/home/me/arc/arcfsmessages`.

- Make a collection:

```

$ mkdir -p mnt/home/me
$ stat mnt/home/me
  File: 'mnt/home/me'
  Size: 0          Blocks: 0          IO Block: 4096   directory
Device: 19h/25d Inode: 3              Links: 2
Access: (0755/drwxr-xr-x)  Uid: ( 500/  me)   Gid: ( 100/  users)
Access: 1970-01-01 01:00:00.000000000 +0100
Modify: 1970-01-01 01:00:00.000000000 +0100
Change: 1970-01-01 01:00:00.000000000 +0100

```

Note that so far there is no security handling implemented, so that all files have mode 0644, all collections have mode 0755, everything is owned by the user that mounts the system.

- Create some entries in your collection:

```
$ cd mnt/home/me
$ emacs -nw fish
catfish
$ cat fish
catfish
$ cp fish fish2
$ ls -la
drwxr-xr-x 5 me users 0 1970-01-01 01:00 .
drwxr-xr-x 3 me users 0 1970-01-01 01:00 ..
-rw-r--r-- 1 me users 8 1970-01-01 01:00 fish
-rw-r--r-- 1 me users 28 1970-01-01 01:00 fish~
-rw-r--r-- 1 me users 8 1970-01-01 01:00 fish2
```

- You probably want to remove that annoying emacs backup file

```
$ rm fish~
$ ls
fish  fish2
```

- Maybe you want fish2 in a separate collection

```
$ mkdir sea_creatures
$ mv fish2 sea_creatures
```

- Maybe you're not happy with the collection name

```
$ mv sea_creatures creatures
$ ls creatures
fish2
```

- Moving the collection out of ARC storage:

```
$ cd ../../..
$ mv mnt/home/me/creatures .
```

- The collection is no longer in the ARC storage:

```
$ find mnt/home/me/
mnt/home/me/
mnt/home/me/fish
```

- Unmounting ARCFs; To unmount, refer to fuse documentation. Usually you can do

```
$ fusermount -u mnt
```

If this for some reason this doesn't work, you can do

```
$ sudo umount -f mnt
```

or even

```
$ sudo umount -l mnt
```

Now, remounting the system, everything should still be there:

```
$ python arcfs.py ./mnt

$ find mnt/home/me/
mnt/home/me/
mnt/home/me/fish

$ fusermount -u mnt
```