



---

NORDUGRID-XXXXXXX-NN

24/7/2008

## DOCUMENTATION OF THE ARC1 STORAGE SYSTEM

*First prototype status and plans*

Zsombor Nagy\*

---

\*zsombor@niif.hu



# Contents

<b>1</b>	<b>Design Overview</b>	<b>5</b>
1.1	IDs used in the system . . . . .	5
1.2	GUIDs and Logical Names (LNs) . . . . .	5
1.3	The Bartenders . . . . .	7
1.4	The Librarians . . . . .	7
1.5	The A-Hash . . . . .	8
1.6	The Shepherds . . . . .	8
1.7	Heartbeats and replication . . . . .	8
<b>2</b>	<b>Use cases</b>	<b>9</b>
2.1	Downloading a file . . . . .	9
2.2	Uploading a file . . . . .	10
2.3	Removing a file . . . . .	12
<b>3</b>	<b>Technical description and implementation status</b>	<b>13</b>
3.1	A-Hash . . . . .	14
3.1.1	Functionality . . . . .	14
3.1.2	Prototype status and plans . . . . .	14
3.1.3	Data model . . . . .	14
3.1.4	Interface . . . . .	14
3.2	Librarians . . . . .	15
3.2.1	Functionality . . . . .	15
3.2.2	Prototype status and plans . . . . .	15
3.2.3	Data model . . . . .	15
3.2.4	Interface . . . . .	17
3.3	Shepherds . . . . .	18
3.3.1	Functionality . . . . .	18
3.3.2	Prototype status and plans . . . . .	18
3.3.3	Data model . . . . .	18
3.3.4	Interface . . . . .	19
3.3.5	Backend modules . . . . .	19
3.4	Bartenders . . . . .	21
3.4.1	Functionality . . . . .	21

3.4.2	Prototype status and plans . . . . .	21
3.4.3	Data model . . . . .	21
3.4.4	Interface . . . . .	21
3.5	Client tools . . . . .	23
3.6	Other plans . . . . .	24
3.6.1	Integrating third-party storage solutions . . . . .	24
3.6.2	Security . . . . .	24

# Chapter 1

## Design Overview

The ARC1 storage system is a distributed system for storing replicated *files* on several file storage nodes and manage them in a global namespace. The files can be grouped into *collections*, and a collection can contain sub-collections and sub-sub-collections in any depth. There is a dedicated *root collection* to gather all collections to the global namespace. This hierarchy of collections and files can be referenced using *Logical Names (LNs)*. The users can use this global namespace as they were using a local filesystem. Files can be transferred by multiple transfer protocols, and the client side tools hide this from the user. The replicas of the files are stored on different storage nodes. A storage node should have one or more data transfer service running (e.g. HTTP(S), FTP(S), GridFTP, ByteIO<sup>1</sup>, etc.), and one of the services of the ARC1 storage system is needed to manage each storage node and to integrate it into the system. And there are plans to provide a way to integrate third-party storage solutions into the namespace of the ARC1 storage system. The main services are shown on Figure 1.1.

### 1.1 IDs used in the system

There are a number of IDs used in the ARC storage system, such as:

- Each file and collection has a globally unique ID called **GUID**.
- The files and collections are organized into a global hierarchical namespace and can be referred to using paths of this namespace called **Logical Names (LNs)**.
- Each service has a unique **serviceID** which can be used to get an endpoint reference from the information system. We need an endpoint reference which is an address (URL) which we could connect to.
- The Shepherds in the system identify their files using a **referenceID**.
- The **location** of a replica consists of two IDs: the ID of the Shepherd and the ID of the file within the Shepherd: (serviceID, referenceID).

### 1.2 GUIDs and Logical Names (LNs)

The syntax of Logical Names is `/[path]` or `<GUID>[/<path>]` where [...] indicates optional parts.

Each file and collection has a GUID which is globally unique, so they can be unambiguously referred using this GUID, that's why a single GUID is a Logical Name itself.

In a collection each entry has a name, and this entry can be a sub-collection, in which there are files and sub-sub-collections, etc.

Example on Figure 1.2: if we have a collection with GUID 1234, and there is a collection called **green** in it, and in **green** there is another collection called **orange**, and in **orange** there is a file called **huge**, then we

---

<sup>1</sup>OGSA ByteIO Working Group (BYTEIO-WG), <https://forge.gridforum.org/projects/byteio-wg/>




Figure 1.1: The components of the ARC1 storage; *A-H*: the **A-Hash**, which is a currently centralized, later distributed database; *L*: the **Librarian**, which stores the metadata and hierarchy of collections and files, the location of replicas, and health data of the Shepherd services, using the A-Hash as database; *B*: the **Bartender**, which provides a high-level interface for the users and for other services; *S*: the **Shepherd**, which provides a simple interface for storing files on storage nodes.

can refer to this file as `/green/orange/huge` if we know the GUID of the starting collection, so let's prefix the path with it: `1234/green/orange/huge`. This is the Logical Name of that file. If there is a well-known system-wide root collection (its GUID could be e.g. 0), then if a LN starts with no GUID prefix, it is implicitly prefixed with the GUID of this well-known root collection, e.g. `/why/blue` means `0/why/blue`.

If a client wants to find the file called `/why/blue`, the system knows where to start the search: the GUID of the root collection. The root collection knows the GUID of `why`, and the (sub-)collection `why` knows the GUID of `blue`. If the GUID of this file is 5678, and somebody makes another entry in collection `/why` (= `0/why`) with name `red` and GUID 5678, then the `/why/red` LN points to the same file as `/why/blue`, so it's a hard link. To count the references (the number of hard links) of a file or collection, the GUID of the parent collection(s) is stored as a metadata for each file and collection.

Each VO<sup>2</sup> should create a VO-wide root collection, and put it in the generic root collection, e.g. if a VO called `vo1` creates a collection called `vo1` as a sub-collection of the root collection (which has the GUID 0), then it can be referred as `0/vo1` or just `/vo1`. Then this VO can create some files, and put them in this `/vo1` collection, e.g. `/vo1/file1`, etc. Or sub-collections, e.g. `/vo1/col1`, `/vo1/col2/file3`, etc. For this the VO does not need to install any service. These files and collections can be created using a Bartender service.

<sup>2</sup>Virtual Organization, [http://www.globus.org/grid\\_software/security/voms.php](http://www.globus.org/grid_software/security/voms.php)

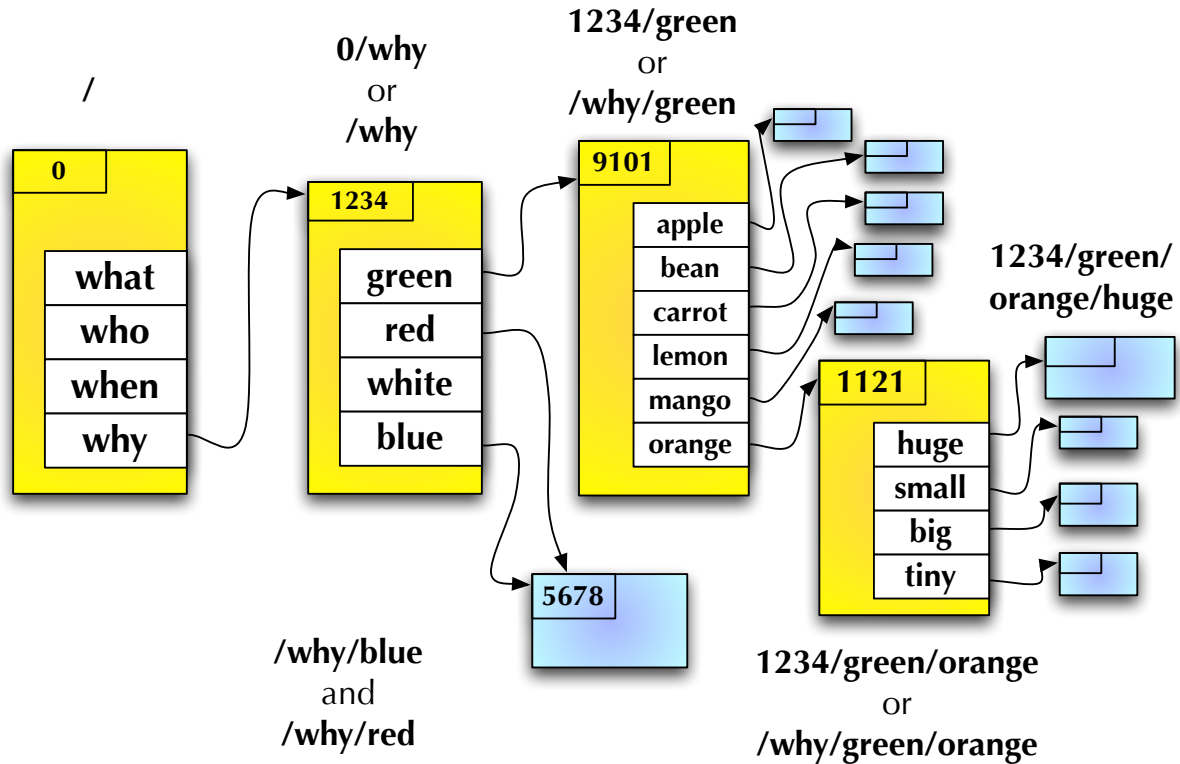


Figure 1.2: Example of the hierarchy of the global namespace

### 1.3 The Bartenders

Clients can access the storage system through a Bartender service. If a client wants to create a collection, upload or download a file, the first step is to connect to a Bartender. The Bartender then asks a Librarian to resolve Logical Names and return metadata, then initiates file transfers on some storage node by asking the Shepherd service of that node, then returns the transfer URL (*TURL*) to the client, which allows the client to directly transfer from/to the storage node. So the data transfer itself is not going through the Bartender, it is performed over a direct link between a storage node and the client.

There could be any number of independent Bartender services in the system which provides high-availability and load-balancing.

### 1.4 The Librarians

The Librarian is capable of managing the hierarchy and metadata of files and collections, and health information of the Shepherd services. Each file and collection in the Librarian has a globally unique ID (GUID). A collection contains files and other collections, and each of these entries has a name unique within the collection very much like entries in a usual directory on a local filesystem. Besides files and collections the Librarian stores a third type of entries called Mount Points which are references to external services which will be in the future used to mount the namespace of third-party storages to our global namespace and make the files on a third-party storage available through the interface of the ARC storage system.

The Librarian also manages information about registered Shepherd services which are associated with a storage node, and receives heartbeat messages from them and change replica states automatically if needed.

The Librarian uses the A-Hash as database, that's why there could be any number of independent Librarian services (all using the same A-Hash) which provides high-availability and load-balancing.

## 1.5 The A-Hash

The A-Hash is a currently centralized, but later distributed service capable of consistently storing objects containing property-value pairs organized in sections. All metadata about files and collections are stored in the A-Hash, and some information about Shepherd services is stored in it as well. The A-Hash itself does not interpret the data, it just stores tuples of strings.

## 1.6 The Shepherds

When a new file is put into the system the number of needed replicas is given for the file. The file replicas are stored on different storage nodes, for each storage node there is Shepherd service which manages the storage node, reports its health state to the Librarian and provides the interface for initiating file transfer.

The file-naming used by a Shepherd has nothing to do with the hierarchy of collections, or Logical Names. When a new replica upload is initiated, the Shepherd generates an ID which refers to it within that Shepherd. Each Shepherd has a unique ID itself, so with these IDs the replica can be unambiguously referenced, this is called a *Location*. The namespace of these Locations has nothing to do with the namespace of GUIDs or the namespace of Logical Names. It consists of two IDs: the ID of the Shepherd and the ID of the file within the Shepherd: (**serviceID**, **referenceID**)

## 1.7 Heartbeats and replication

Each Shepherd should periodically send heartbeats to a Librarian service with information about replicas whose state changed since the last heartbeat, the Librarian stores these file lists (which contains the GUIDs of the files as well), and if it doesn't receive a heartbeat for a Shepherd in a given time, it invalidates all the replicas the Shepherd stores. This invalidating means that the state of that location will be **offline**.

If a Shepherd finds out that a file is missing or has a bad checksum, it reports that the file is **invalid** to the Librarian immediately, and the Librarian alters the state of the given replica of the file. The Shepherd tries to recover its replica by downloading it from another Shepherd. In order to do this the Shepherd contacts a Bartender and asks for the file. The Bartender chooses a valid replica, initiates file transfer by a Shepherd having a valid replica, and returns the TURL to the Shepherd with the invalid replica. The Shepherd with the invalid replica downloads the file from the other Shepherd, and if everything is OK, signals to the Librarian that the replica is **alive** again.

The Shepherds periodically ask the Librarian whether the files they store have enough replicas. If a Shepherd finds that one of the files has not enough replica it turns to a Bartender offering replication. The Bartender chooses a Shepherd, initiates a put request then returns the TURL to the offering Shepherd which could upload the replica. The Shepherd who now has the new replica notifies the Librarian that the file is **alive**. The Librarian sets the state of this new replica.



## Chapter 2

## Use cases

### 2.1 Downloading a file

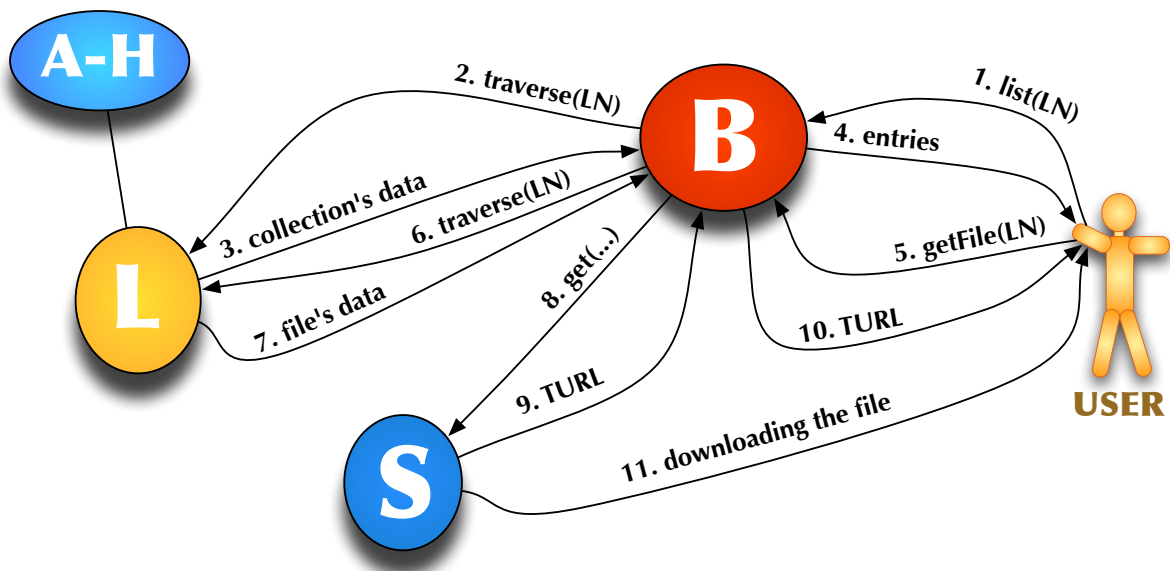


Figure 2.1: Downloading a file

We want to download a file about which we know that it is somewhere in our home collection on the storage (see Figure 2.1). The LN of our home collection is e.g. `/ourvo/users/we` (here a ‘home collection’ could be just a collection which is given to us by our VO). We can get a list of entries in this collection from any Bartender.

1. We need to find a Bartender. Maybe we have a cached list of recently used Bartenders or we can get one from the information system. When we have an endpoint reference of a Bartender, we could call its list method with the LN `/ourvo/users/we`.
2. The Bartender has to find a Librarian service, again using its cache of recently used Librarian services or get a new one from the information system. When the Bartender has an endpoint reference of a Librarian service, it could ask the Librarian to traverse the LN `/ourvo/users/we`.
3. The Librarian needs the A-Hash service to access the stored data, when it has the endpoint reference of the A-Hash service, it could get the information about the root collection, which contains the GUID of the `ourvo` sub-collection. Then the Librarian gets the entries of this `ourvo` collection, and in it it can find the GUID of `users`, and in the entries of `users` there is the GUID of `we`, which the Librarian returns to the Bartender with all the metadata.

4. The Bartender now has the GUID and the metadata of the collection `/ourvo/users/we`, including the list of its entries. This is returned to us.
5. So we get the list of our `/ourvo/users/we` collection, and now we realize that the file we want has the LN `/ourvo/users/we/thefilewewant` and we know the GUID of it as well: e.g. `a4b2e`. (Of course we know the GUID of the `/ourvo/users/we` collection too, which is e.g. `13245` and using this we could refer to our file as `13245/thefilewewant` which means the entry called `thefilewewant` in the collection with a GUID `13245`.) We connect a Bartender again (the same one or maybe another one) to get the file with any of these LNs, the `a4b2e` is the fastest solution because the Bartender need not to look up the whole LN again in the Librarian, a well-written client API should use this. With the get request we give the Bartender the list of transfer protocols we are able to use.
6. The Bartender contacts the Librarian to get the locations of the replicas of this file.
7. The Librarian returns all the metadata of the requested LN.
8. The Bartender chooses a 'location' which consists of the ID of a Shepherd, and the referenceID of the file within the Shepherd. Using the information system or its local cache it could get the endpoint reference of the Shepherd. The Bartender initiates a transfer by the Shepherd, if the Shepherd supports one of the transfer protocols we give, it can create a transfer URL (TURL) with a protocol we can download.
9. The Shepherd returns the TURL to the Bartender.
10. The Bartender returns the TURL to us.
11. Now we have a TURL from which we can download it.

## 2.2 Uploading a file

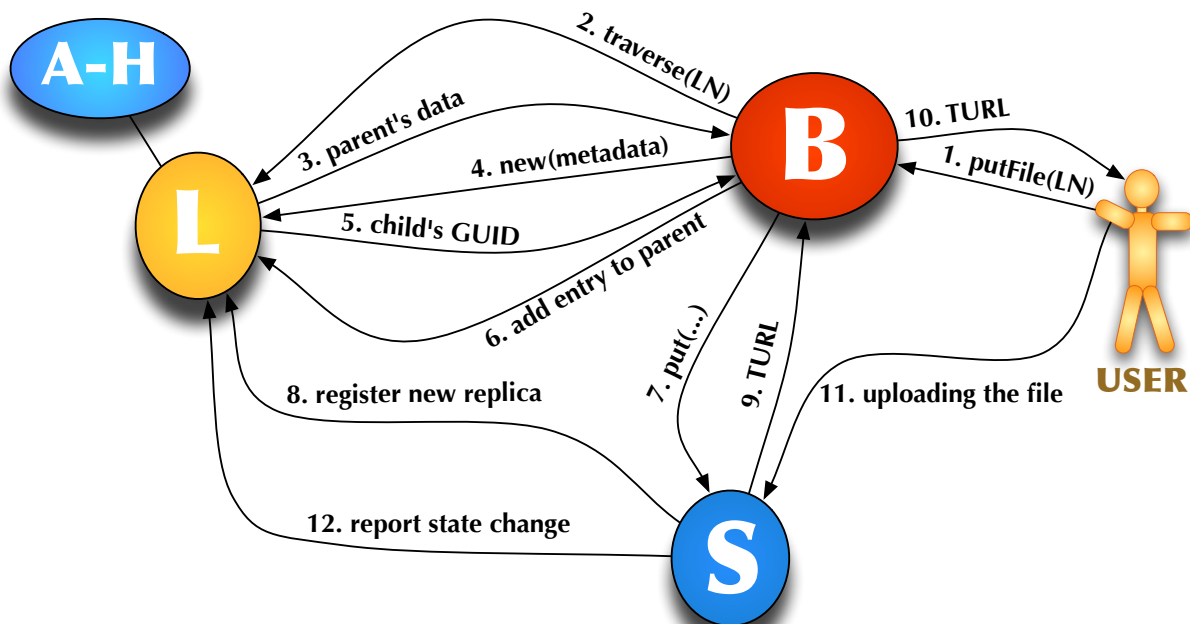


Figure 2.2: Uploading a file

We have a file on our local disk we want to upload to a collection called `/ourvo/common/docs`. (See Figure 2.2.)

1. We contact a Bartender to put the file, we give the size and checksum and other metadata, and the transfer protocols we want to use. And of course we give the Logical Name we want to be the name of the file, which in this case will be `/ourvo/common/docs/proposal.pdf`
2. The Bartender ask a Librarian to traverse this LN.
3. If the Librarian can traverse the whole LN then this LN is already exists, but if the LN is still available and the parent exists, the Librarians response contains the GUID and metadata of the parent collection.
4. Then the Bartender creates a new file entry within the Librarian with all the information we gave.
5. The Librarian returns the GUID of this new entry.
6. Then the Bartender add the name `proposal.pdf` and the new GUID to the collection `/ourvo/common/docs` and from now on there will be a valid LN `/ourvo/common/docs/proposal.pdf` which points to a file which has no replica at all. If someone tried to download the file called `/ourvo/common/docs/proposal.pdf` now, would get an error message 'try again later'.
7. The Bartender (using the information system) chooses a Shepherd and gets its endpoint reference. Then the Bartender initiates uploading of the file to the Shepherd: the request includes the size and checksum of the file, the GUID, and the protocols we are able to use.
8. The Shepherd creates a transfer URL and a referenceID for this file and registers the GUID of the file in its own database and reports to the Librarian that there is a new replica with state **creating**. The Librarian gets the message from the Shepherd and creates a new entry in the locations list of the given file with the serviceID and the referenceID the Shepherd have just reported. If someone tries to download this file now, still gets a 'try again later' error message, because this new replica is still not **alive**.
9. The Shepherd returns the the TURL to the Bartender.
10. The Bartender returns the TURL to us.
11. Then we can upload the file to this TURL.
12. The Shepherd detects that the file is arrived and reports the change of state to **alive** to the Librarian who alters the state in the given file-entry. At this point the file has only one replica.
  - The Shepherd periodically checks the Librarian if this is less than the needed replica number, and if it is then it initiates creating a new replica by a Bartender.
  - The Bartender chooses another Shepherd, initiates the transfer then returns the TURL to the first Shepherd which uploads the file to the new Shepherd.
  - All the Shepherds check periodically whether their files have enough replica, and if any of them find that there is more replica needed, it initiates creating a new. If more than one Shepherd of course could cause that there will be more replicas than needed. If a Shepherd finds out that a file has more replicas than needed it notifies a Bartender about it.
  - The Bartender ask the Librarian about all Shepherds this file has replicas on, flags this file as 'removing a replica' which prevents other Bartenders to remove an other replica accidentally, then make a decision of which one is to be removed, then contacts the chosen Shepherd and asks it to remove the replica. The Shepherd then notify the Librarian, and the Librarian removes the replica, and removes the flag 'removing a replica' as well.
  - If we cannot upload the file to the given TURL for some reason, we should remove the file entry from the collection, or we should call **addReplica** to get a new TURL without removing and recreating the file.

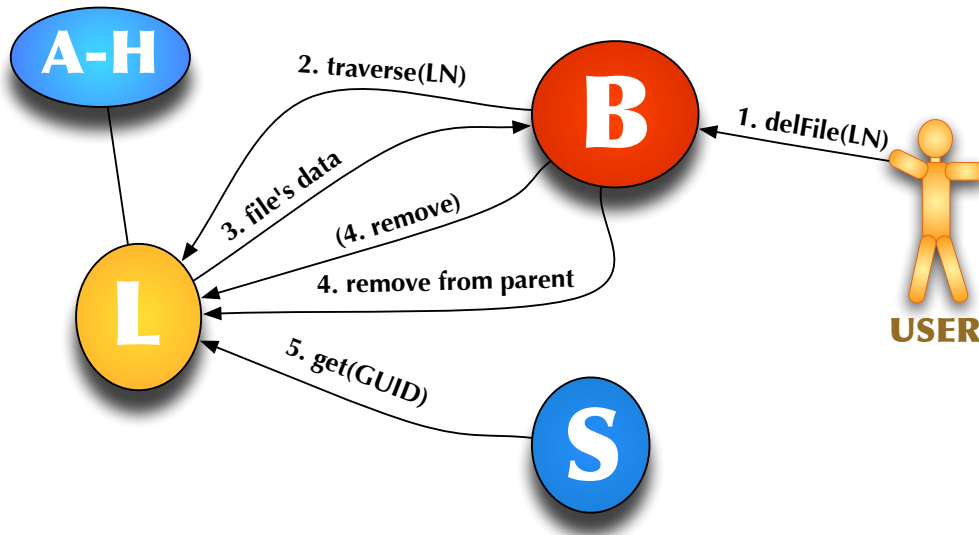


Figure 2.3: Removing a file

## 2.3 Removing a file

1. If we want to remove a file, we should connect to a Bartender with the LN of the file we want to remove.
2. The Bartender asks the Librarian to traverse the LN and return the list of parent collections of this file.
3. The Librarian returns the data.
4. If the file has only parent, it asks the Librarian to remove the file entry itself, and the entry from the parent collection. If the file has more parent collections (hard links) the Bartender only removes the entry from the parent collection and removes the parent from the list of parent collections of this file.
5. Next time the Shepherd does its periodic check, it asks the Librarian about each of its stored replicas, and finds out that one of them no longer exists, so it removes the replica from the storage node.

## Chapter 3

# Technical description and implementation status

The services are written in Python and running in the HED<sup>1</sup> hosting environment. The HED itself is written in C++, but there are language bindings which allow us to write services in other languages, e.g. in Python or Java. The source tree of the storage services are in the NorduGrid Subversion repository<sup>2</sup>.

The current version of the prototype has no information system and no security, these are soon to be integrated to the system.

The information system is needed to discover services, and to translate *serviceIDs* to endpoint references (URLs). Currently the URLs are written in the configuration files, and the Shepherd services are reporting their URLs to the Librarian, so the Bartender could ask for all alive Shepherds.

The security is needed to do proper authorization of the users, and to manage access policies of files and collections. ARC has its own policy language, for each file and collection there will be a policy XML document stored as a metadata. The storage services will use the properties extracted from the communication channel and these policies to make authorization decisions. If the properties and the policies are present, the decision is actually made by the security framework of HED.

Further prototype statuses and plans can be found below within each section about the services.

---

<sup>1</sup>The ARC container - <https://www.knowarc.eu/documents/Knowarc.D1.2-2.07.pdf>

<sup>2</sup><http://svn.nordugrid.org/trac/nordugrid/browser/arc1/trunk/src/services/storage>

## 3.1 A-Hash

### 3.1.1 Functionality

The A-Hash is a distributed service capable of storing objects containing property-value pairs grouped in sections in a scalable manner. Each object has an arbitrary string ID, and contains any number of property-value pairs grouped in sections, where property, value and section are arbitrary strings. There could only be one value for a property in a section.

If you have an ID, you can get all property-value pairs of the corresponding object with the *get* method, or you could specify only which sections or properties you need. You can add or remove property-value pairs of an object or delete all occurrences of a property or create a new object with the *change* method, and you can specify conditions, which means the change is only applied if the given conditions are met.

### 3.1.2 Prototype status and plans

The A-Hash service currently implemented as a single central service, which stores the data on disk in separate files per object. In the fall of 2008 it will be reimplemented on a distributed hash table (DHT) algorithm, for example the Chord algorithm with a consistency solution called Etna on top of it. This reimplementation hopefully won't change the interface of the service.

### 3.1.3 Data model

- *ID* is an arbitrary string
- *object* contains property-value pairs in sections, technically it is a list of key-value pairs where the key is a (section, property) tuple

### 3.1.4 Interface

**get(ids, neededMetadata)** returns *getResponse* which is a list of (*ID*, *object*) pairs.

The *ids* is a list of string *IDs*, *neededMetadata* is a list of (*section*, *property*) pairs. For each *ID* it returns all the *values* for each *property* in each *section* (filtered by *neededMetadata*), so *object* is a list of (*section*, *property*, *value*) tuples.

**change(changeRequest)** returns *changeResponse* which is a list of (*changeID*, *success*, *failedConditionID*) tuples.

*changeRequest* is a tuple of (*changeID*, *ID*, *changeType*, *section*, *property*, *value*, *conditions*), where *changeID* is an arbitrary ID to identify in the response which change was successful; *ID* points to the object we want to change; *changeType* can be 'set' (set the property within the section to value), 'unset' (remove the property from the section regardless of the value), 'delete' (removes the whole object), *conditions* is a list of (*conditionID*, *type*, *section*, *property*, *value*) tuples, where *type* could be 'is' (the property in the section is set to the value), 'isnot' (the property in the section is not set to the value), 'isset' (the property of the section is set to any value), 'unset' (the property of the section is not set at all). If all conditions are met, tries to apply changes to the objects, creates a new object if a previously non-existent ID is given. If one of the conditions is not met, returns the ID of the failed condition.

## 3.2 Librarians

### 3.2.1 Functionality

The Librarian manages a tree-hierarchy of files, grouping them into collections. There is a root collection with a well-known GUID which can be used as starting point when resolving Logical Names. If you create a new collection with the method *new*, the Librarian generates a new GUID, but does not insert it into the tree-hierarchy which can be done by adding this GUID as a new entry to one of the existing collection using the *modifyMetadata* method of the existing collection which makes it the parent of the new collection. A collection can be closed via metadata modification which cannot be undone and prevents files to be added or removed from this collection. A new file also can be created with the *new* method which returns the newly generated GUID of the new file entry which should be added to a parent collection to insert it into the global namespace. A file has a list of locations where its replicas are stored, this list too can be manipulated with *modifyMetadata*. The access policies of the files and collections are also stored as metadata. The *remove* method deletes an entry from the Librarian. The *traverseLN* method try to traverse Logical Names by walking the hierarchy of the namespace and to return the GUID of the entry pointed by the LN. After you have a GUID of file, collection or mount point, you can get all the information using the *get* method.

### 3.2.2 Prototype status and plans

The Librarian service currently implements all the methods below, but doesnt do very much error checking. This should be changed, the Librarian should check the validity of metadata, and forbid some cases, e.g. reopen a closed collection.

### 3.2.3 Data model

Each librarian entry has a unique ID called *GUID*.

The Librarian uses the A-Hash to store all the data about the files and collections. The A-Hash is capable of storing property-value pairs organized in sections, which actually means that it stores (*section*, *property*, *value*) tuples where each member is simply a string, e.g. ('entry', 'type', 'collection') or ('ACL', 'john-smith', 'owner') or ('timestamps', 'created', '1196265901') or ('locations', '64CDF45F-DDFA-4C1D-8D08-BCF7810CB2AB 9A293F27DC86', 'sentenced'). There could be only one *value* for a (*section*, *property*) pair.

- A **collection** is a list of files and other collections, which are in parent-children relationships forming a tree-hierarchy. Each entry has a name which is only valid within this collection, and it is unique within the collection. Each entry is referenced by its GUID. So the metadata sections of a collection are as follows:

**entry** section

- *type*: 'collection'

**entries** section

- (*name*, *GUID*) *pairs*: a Collection is basically a list of name-GUID pairs.

**timestamps** section

- *created*: timestamp of creation
- *modified*: timestamp of last modification

**states** section

- *closed*: if the collection is closed, then nothing can be added to its contents

**policies** section

- XML representations of access policies

**metadata** section

- any other arbitrary metadata

- A **file** entry contains the following sections:

**entry** section

- *type*: ‘file’

**locations** section

- (*location, state*) pairs, where a location is a (*serviceID, referenceID*) pair serialized as a string, where *serviceID* is the ID of the Shepherd service storing this replica, *referenceID* is the ID of the file within that Shepherd service, and state could be ‘**alive**’ (if the replica passed the checksum test, and the storage element storing it is healthy), ‘**invalid**’ (if the replica has wrong checksum, or the storage element claims it has no such file), ‘**offline**’ (if the storage element is not reachable, but may have a valid replica), ‘**creating**’ (if the replica is in the state of uploading), ‘**sentenced**’ (if the replica is marked for deletion)

**timestamps** section

- *created*: timestamp of creation
- *modified*: timestamp of last modification (e.g. modification of metadata)

**states** section

- *size*: the file size in bytes
- *checksum*: checksum of the file
- *checksumType*: the name of the checksum method
- *neededReplicas*: how many valid replicas should this file have

**policies** section

- XML representations of access policies

**metadata** section

- any other arbitrary metadata

- There is one more type of Librarian entries called **mount point** which is a reference to a service which is capable of handling a subtree of the namespace. The properties of a mount point in sections:

**entry** section

- *type*: ‘mountpoint’

**mount** section

- *target*: the ID of the service
- *ID*: an ID within the service (optional)

**timestamps** section

- *created*: timestamp of creation
- *modified*: timestamp of last modification (e.g. modification of metadata)

**policies** section

- XML representations of access policies

**metadata** section

- any other arbitrary metadata

- The Librarian stores information about the Shepherds, so each Shepherd has a GUID as well. There is an entry (with GUID ‘1’ by default) which contains the GUID and the timestamp of the last heartbeat for each registered Shepherd:

**nextHeartBeat** section

- (ID, timestamp) pairs

**serviceGUID** section

- (ID, GUID) pairs

- For each Shepherd there is a separate entry with the list of files:

**entry** section

- *type*: ‘shepherd’

**file** section

- (*referenceID, GUID*) pairs for each replica stored on the Shepherd



### 3.2.4 Interface

**new(newRequestList)** returns a list of (*requestID*, *GUID*, *success*)

*newRequestList* is a list of (*requestID*, *metadata*) where *requestID* is an arbitrary ID used to identify this request in the list of responses; *metadata* is a list of (*section*, *property*, *value*) tuples. This method generates a *GUID* for each request, and inserts the new entry into the A-Hash, then returns the GUIDs of the newly created entries.

**modifyMetadata(modifyMetadataRequestList)** returns a list of (*changeID*, *success*)

*modifyMetadataRequestList* is a list of (*changeID*, *GUID*, *changeType*, *section*, *property*, *value*) where *changeType* can be ‘**set**’ (set the property in the section to the value), ‘**unset**’ (remove the property-value pair from a section), ‘**add**’ (set the property in the section to the value only if it is not exists already).

**get(GUIDs, neededMetadata)** returns *getResponse*

*GUIDs* is a list of GUIDs, *neededMetadata* is a list of (*section*, *property*) pairs indicating only which properties we need, *getResponse* is a list of (*GUID*, *metadata*) where *metadata* is a list of (*section*, *property*, *value*) tuples.

**remove(removeRequestList)** returns a list of (*requestID*, *success*) pairs

*removeRequestList* is a list of (*requestID*, *GUID*) pairs. *success* could be ‘**removed**’ or ‘**failed**: reason’.

**traverseLN(traverseRequestList)** returns *traverseResponseList*

*traverseRequestList* is a list of (*requestID*, *LN*) with the Logical Names to be traversed

*traverseResponseList* is a list of (*requestID*, *metadata*, *GUID*, *traversedLN*, *restLN*, *wasComplete*, *traversedList*) where:

**metadata** is all the metadata of the of traversedLN in the form of (*section*, *property*, *value*) tuples

**GUID** is the *GUID* of the *traversedLN*

**traversedLN** is the part of the *LN* which was traversed, if *wasComplete* is true, this should be the full *LN*

**restLN** is the postfix of the *LN* which was not traversed for some reason, if *wasComplete* is true, this should be an empty string

**wasComplete** indicates whether the full *LN* was traversed

**traversedList** is a list of (*LNpart*, *GUID*) pairs, where *LNpart* is a part of the *LN*, *GUID* is the GUID of the Librarian-entry referenced by that part of the *LN*, the first element of this list is the shortest prefix of the *LN*, the last element is the *LN* without its last part

**report(serviceID, filelist)** returns in *nextReportTime* a number of seconds, which is the timeframe within the Librarian expects the next heartbeat from the Shepherd

*filelist* is a list of (*GUID*, *referenceID*, *state*) tuples containing the state of changed or new files, where *state* could be ‘**invalid**’ (if the periodic self-check of the Shepherd found a non-matching checksum or missing file), ‘**creating**’ (if this is a new file not uploaded yet) or ‘**alive**’ (if the file is uploaded and the checksum is OK).

### 3.3 Shepherds

#### 3.3.1 Functionality

The Shepherd service is capable of managing a storage node. It keeps track all the files it stores with their GUIDs and checksums. The Shepherds periodically send reports to a Librarian indicating that the Shepherd is up and running, and whether some file's state has been changed. The Shepherd periodically checks each file to detect corruption. If a file goes missing or has a bad checksum the Shepherd notify the Librarian about the error referring the file with its GUID. The Shepherd periodically asks the Librarian how many replicas its files have, and if a file has fewer replicas than needed, the Shepherd offers its copy for replication by calling the Bartender.

A Shepherd service is always connected to a file transfer service (e.g. 'HTTP(S)', 'FTP(S)', 'ByteIO', 'GridFTP', etc.). For each supported file transfer service we need a backend module which makes the Shepherd able to communicate with the file transfer service to initiate file transfers, to detect whether a transfer was successful or not, to generate local IDs and checksum, etc.

A file in a Shepherd could be identified with a *referenceID* which is unique within that Shepherd. If we know the *location* of a file, which is the ID of the Shepherd service (*serviceID*) and the *referenceID*, we could get the endpoint reference (URL) of the Shepherd from the information system, then we could call its *get* method with the *referenceID* and a list of transfer protocols we are able to handle (e.g. 'HTTP', 'FTP'), the Shepherd chooses a protocol from this list which it can provide, and create a transfer URL (*TURL*) and returns it along with the *checksum* of the file. We could download the file from this *TURL*, and verify it with the *checksum*. An end user of the storage system does not need to call this *get* method, because the Bartender service will do it, the user just ask the Bartender and get the *TURL*.

Storing a file starts with initiating the transfer with the *put* method of the Shepherd, we should give the *size* and *checksum* of the file and its *GUID* as well. We also specify a list of transfer protocols we are able to use, and the Shepherd chooses a *protocol*, creates a *TURL* for uploading and generates a *referenceID*, then we can upload the file to the *TURL*. Again, the end user just asks the Bartender, and get the *TURL*, the user does not need to call the *put* method of the Shepherd directly.

These *TURLs* are one-time URLs which means that after the client uploads or downloads the file these *TURLs* cannot be used again to do the same. If we want to download the same file twice, we have to initiate the transfer twice, and will get two different *TURLs*.

With the *stat* method we can get some information about a replica, e.g. checksum, GUID, state ('creating', 'alive' or 'invalid'), etc. The *delete* method removes the file.

In normal operation the *put* and *get* calls is made by a Bartender but the actual uploading and downloading is done by the users client. In case of replication a Shepherd with a valid replica initiates the replication, this Shepherd asks the Bartender to choose a new Shepherd, the Bartender initiates putting the new replica on a chosen Shepherd and receives the *TURL*, then the Bartender returns the *TURL* to the initiator Shepherd, which uploads its replica to the given *TURL*.

#### 3.3.2 Prototype status and plans

The current implementation of the Shepherd service has a working *get*, *put*, *stat*, *delete* methods, and a method called *toggleReport* which can be used to simulate storage node failure with the Shepherd not reporting to a Librarian. There is a separate service which provide a subset of the ByteIO interface, and there is an other separate service which is a basic HTTP server, these are both could be used as file transfer services, both have its backend module for the Shepherd. Currently both file transfer services have the problem of using too much memory while transferring files. Further plans include better file transfer services and backend modules for third-party file transfer services.

#### 3.3.3 Data model

A file of a Shepherd service is referenced by its *referenceID*. Each file has a *state* which could be 'creating' when the transfer is initiated but the file is not uploaded yet, 'alive' if the file is uploaded and has a proper

checksum, or **'invalid'** if it does not exist anymore or has a bad checksum. Each file has a *localID* which is used in the backend modules..

### 3.3.4 Interface

**get(getRequestList)** returns list of (*requestID*, *getResponseData*)

*getRequestList* is a list of (*requestID*, *getRequestData*) where *requestID* is an arbitrary ID used in the reply

*getRequestData* is a list of (*property*, *value*) pairs, where mandatory properties are: **'referenceID'** which refers to the file to get and **'protocol'** indicates a protocol the client is able to use, there could be multiple protocols in *getRequestData*.

*getResponseData* is a list of (*property*, *value*) pairs, such as: **'TURL'** is a transfer URL which can be used by the client to download the file; **'protocol'** is the protocol of the TURL; **'checksum'** is the checksum of the replica; **'checksumType'** is the name of the checksum method and **'error'** could contain an error message if there is one.

**put(putRequestList)** returns a list of (*requestID*, *putResponseData*)

*putRequestList* is a list of (*requestID*, *putRequestData*) where *requestID* is an ID used for the response

*putRequestData* is a list of (*property*, *value*) pairs such as **'GUID'**, **'checksum'**, **'checksumType'**, **'size'** (the size of the file in bytes), **'protocol'** (a protocol the client can use, can be multiple) and **'acl'** (for additional access policy).

*putResponseData* is a list of (*property*, *value*) pairs such as: **'TURL'** is the transfer URL where the client can upload the file, **'protocol'** is the chosen protocol of the TURL and **'referenceID'** is the generated ID for this new replica, **'error'** could contain an error message.

**delete(deleteRequestList)** returns a list of (*requestID*, *status*)

*deleteRequestList* is a list of (*requestID*, *referenceID*) pairs selecting the files to remove. The status could be **'deleted'** or **'nosuchfile'**.

**stat(statRequestList)** returns a list of (*requestID*, *referenceID*, *state*, *checksumType*, *checksum*, *acl*, *size*, *GUID*, *localID*)

*statRequestList* is a list of (*requestID*, *referenceID*) where *referenceID* points to the file whose data we want to get. The method returns all the data the Shepherd know about the replica.

### 3.3.5 Backend modules

The Shepherd could communicate with the file transfer services via backend modules. Currently there are two kinds of backend modules, one for the byteio service (which is a simple implementation of a subset of the ByteIO interface) and one for the httpd service (which is a simple HED-based HTTP server).

In both cases the Shepherd service uses two separate directories: one for storing all the files (e.g. **./store**) and one for the file transfers (e.g. **./transfer**). The store directory always contains all the files the Shepherd manages, the transfer directory is empty at the beginning.

Let's see the scenario for the httpd service which should be in a special 'slave' mode for this kind of operation: if a client asks for a file called **file1**, and this file is in the store directory (**./store/file1**), then the Shepherd service creates a hardlink into the transfer directory (e.g. **./transfer/abc**) and sets this file read-only. If the httpd service is configured that way that it handles the HTTP path **/prb** and it is serving files from the directory **./transfer** then after the hardlink is created, we have this URL for this file: **http://localhost:60000/prb/abc**. Now we can give this URL to the client. Then the client **GETs** this URL and gets the file. The httpd service removes (unlinks) this file immediately after the **GET** request arrived, which makes this **http://localhost:60000/prb/abc** URL invalid (so this is a one-time URL), but because of the hardlink the file is still there in the store directory, it is just removed from the transfer directory. Now if some other user wants this file, the Shepherd creates an other hardlink, e.g. **./transfer/qwe** and now we have an URL **http://localhost:60000/prb/qwe**.

If a client wants to upload a new file, then the Shepherd creates an empty file in the store directory, e.g. **./store/file2** and creates a hardlink into the transfer directory, e.g. **./transfer/oiu** and makes it

writable, and now we have a URL `http://localhost:60000/prb/oiu`, and the client is able to do a PUT to this URL. When the client PUTs the file there, the httpd service immediately removes the uploaded file from the transfer directory, but because it has a hardlink in the store directory, the file is stored there as `./store/file2`. The backend module for the httpd service periodically checks whether a new file has two or just one hard links. If it has only one that means that a file is uploaded, so it could notify the Shepherd that the file is arrived. In order to do that, all the backend modules get a callback method ‘file\_arrived’ from the Shepherd.

In case of the byteio service, there is some small differences. The byteio service does not removes the files from the transfer directory, but it calls the backend module via SOAP to notify it that something is happened. The byteio backend has one SOAP method called ‘**notify**’:

**notify(subject, state)** returns ‘OK’ in *notifyResponse*.

When this method is called, the byteio backend module notifies the Shepherd that the file is arrived.

All the backend modules should have this common interface which the Shepherd can use to communicate with the file transfer service:

**prepareToGet(referenceID, localID, protocol)** returns the *TURL*.

Initiate transfer with *protocol* for the file which has these IDs: *localID* and *referenceID*. The reason for including here the *referenceID* as well is that this information could be used by the backend module later, e.g. when the transfer finished and the state of the file needs to be changed.

**prepareToPut(referenceID, localID, protocol)** returns the *TURL*.

Initiate transfer with *protocol* for the file which has these IDs: *localID* and *referenceID*.

**copyTo(localID, turl, protocol)** returns *success*.

Upload the file referenced by *localID* to the given *TURL* with the given *protocol*.

**copyFrom(localID, turl, protocol)** returns *success*.

Download the file from the given *TURL* with the given *protocol*, and store it as *localID*.

**list()** returns a list of *localIDs* currently in the store directory.

**getAvailableSpace()** returns the available disk space in bytes.

**generateLocalID()** returns a new unique *localID*.

**matchProtocols(protocols)** only leave that protocols in the list *protocols* which are supported by this file transfer service.

**checksum(localID, checksumType)** returns the checksum of the file referenced by *localID*, which checksum is generated by the method *checksumType*.

## 3.4 Bartenders

### 3.4.1 Functionality

The Bartender provides an easy to use interface of the ARC1 storage system to the users. You can put, get and delete files using their logical names (*LN*s) with the *putFile*, *getFile* and *delFile* methods, create, remove and list collections with *makeCollection*, *unmakeCollection* and *list*. The metadata of a file or collection (e.g. whether the collection is closed, number of needed replicas, access policies) can be changed with *modify*. A *stat* call gives all the information about a file or collection, and you can move (or hardlink) collections and files within the namespace with *move*. You can upload an entirely new replica to a file (e.g. if the file lost all its replicas, or when a Shepherd service offers its replica for replications) with *addReplica*.

### 3.4.2 Prototype status and plans

The methods mentioned in the above section are all implemented, but need more error-checking and metadata-checking. There are plans of adding new methods, e.g. a *copy* method or a *glob* method for file pattern matching. The current version does not support closed (unmodifiable) collections yet.

### 3.4.3 Data model

The Bartender interface uses mostly Logical Names (*LN*s), which have the syntax of: <GUID>/<path> where both sides can be omitted (and in the case of a sole GUID we don't need the slash either), e.g. `afg342/foo` is an entry called `foo` in the collection with GUID `afg342`; the LN `'f36a7481'` refers to the file or collection with GUID `f36a7481`; `/vo/dir/stg` points to the entry which is reachable from the root collection using the given path; and `/` simply refers to the root collection. The term '*metadata*' here refers to a list of property-value pairs organized in sections, see the data model description in Section 3.2.3.

### 3.4.4 Interface

**putFile(putFileRequest)** returns a list of (*requestID*, *success*, *TURL*, *protocol*)

*putFileRequest* is a list of (*requestID*, *LN*, *metadata*, *protocols*), where *requestID* is an arbitrary ID which will be used in the response; *LN* is the chosen Logical Name of the new file, *protocols* is a list of protocols we can use for uploading, *metadata* is a list of (*section*, *property*, *value*) tuples where properties could be in the '*states*' section: '*size*', '*checksum*', '*checksumType*', and '*neededReplicas*', policy documents in the '*policies*' section and any other property-value pairs in the '*metadata*' section. The returned *TURL* is a URL with a chosen *protocol* to upload the file itself, the *success* string could be '*done*', '*missing metadata*', '*parent does not exists*', '*internal error: reason*', etc.

**getFile(getFileRequest)** returns a list of (*requestID*, *success*, *TURL*, *protocol*)

*getFileRequest* is a list of (*requestID*, *LN*, *protocols*) where *requestID* is used in the response, *LN* is the Logical Name referring to the file we want to get, *protocols* is a list of transfer protocols the client supports. In the response *TURL* is the transfer URL using *protocol*, with which we can download the file, *success* could be '*done*', '*not found*', '*is not a file*', '*file has no valid replica*', '*error while getting TURL: reason*', etc.

**delFile(delFileRequest)** returns a list of (*requestID*, *status*)

*delFileRequest* is a list of (*requestID*, *LN*) with the Logical Name of the file we want to delete. The status in response could be '*deleted*' or '*nosuchLN*'.

**stat(statRequest)** returns a list of (*requestID*, *metadata*)

*statRequest* is a list of (*requestID*, *LN*) with the Logical Name of the file or collection we want to get information about, and it returns *metadata* which is a list of (*section*, *property*, *value*) tuples according to the data model of the Librarian (see Section 3.2.3)

**makeCollection(makeCollectionRequest)** returns a list of (*requestID*, *success*)

*makeCollectionRequest* is a list of (*requestID*, *LN*, *metadata*) where *metadata* is a list of (*section*, *property*, *value*) tuples where in the ‘**entries**’ section there could be the initial content of the catalog in the form of name-GUID pairs (these entries will be hard links to the given GUIDs with the given name), in the ‘**states**’ section there is the ‘**closed**’ property (if it is true then no more files can be added or removed later), in the ‘**policies**’ section there could be some access policies, and in the ‘**metadata**’ section there could be any other metadata in key-value pairs. The *success* in the response could be ‘done’, ‘**LN exists**’, ‘**parent does not exist**’, ‘**failed to create new catalog entry**’, ‘**failed to add child to parent**’, ‘**internal error**’, etc. [*list*(*listRequest*, *neededMetadata*)] returns *listResponse*.

*listRequest* is a list of (*requestID*, *LN*) where *LN* is the Logical Name of the collection (or file) we want to list, *neededMetadata* is a list of (*section*, *property*) pairs which filters the returned metadata.

*listResponse* is a list of (*requestID*, *entries*, *status*) where *entries* is a list of (*name*, *GUID*, *metadata*) where *metadata* is a list of (*section*, *property*, *value*) tuples according to the data model of the Librarian (Section 3.2.3), the *status* could be ‘**found**’, ‘**not found**’, ‘**is a file**’ (because only collections can be listed).

**move(moveRequest)** returns a list of (*requestID*, *status*).

*moveRequest* is a list of (*requestID*, *sourceLN*, *targetLN*, *preserveOriginal*) where *sourceLN* is the Logical Name referring to the file or collection we want to move (or just rename) and *targetLN* is the new path, and if *preserveOriginal* is true the *sourceLN* would not be removed, so with *preserveOriginal* we actually creating a hard link. The status could be ‘**moved**’, ‘**nosuchLN**’, ‘**targetexists**’, ‘**invalidtarget**’, ‘**failed adding child to parent**’, ‘**failed removing child from parent**’

**unmakeCollection(unmakeRequest)** returns a list of (*requestID*, *success*).

*unmakeRequest* is a list of (*requestID*, *LN*) with the Logical Names of the collections we want to remove. *success* could be ‘**removed**’, ‘**no such LN**’, ‘**collection is not empty**’, ‘**failed: reason**’.

**modify(modifyRequest)** returns a list of (*changeID*, *success*) *modifyRequest* is a list of (*changeID*, *LN*, *changeType*, *section*, *property*, *value*) where *changeType* can be ‘**set**’ (set the *property* in the *section* to *value*), ‘**unset**’ (remove the *property-value* pair from the *section*), ‘**add**’ (set the *property* in the *section* to *value* only if it is not exists already). *success* could be ‘**no such LN**’, ‘**set**’, ‘**unset**’, ‘**entry exists**’, ‘**failed: reason**’.

## 3.5 Client tools

In the first prototype release there is one client tool called `arc_storage_cli.py`, which is written in Python, and only need a basic Python installation to run. It is capable of communicating with a given Bartender service, and uploading and downloading TURLs via HTTP.

The methods can be listed with:

```
$ arc_storage_cli.py
Usage:
  arc_storage_cli.py <method> [<arguments>]
Supported methods: stat, make[Collection], unmake[Collection], list, move,
  put[File], get[File], del[File]
```

Without arguments, each method prints its own help:

```
$ arc_storage_cli.py move
Usage: move <sourceLN> <targetLN>
```

Uploading, downloading and stat files:

```
$ cat testfile
This is a testfile.
$ arc_storage_cli.py put testfile /tmp/
- The size of the file is 20 bytes
- The md5 checksum of the file is 9a9dffa22d227afe0f1959f936993a80
- ARC_BARTENDER_URL environment variable not found, using http://localhost:60000/Bartender
- Calling the Bartender's putFile method...
- done in 0.08 seconds.
- Got transfer URL: http://localhost:60000/httpd/d15900f5-34ee-4bba-bb10-73d60d1c0d75
- Uploading from 'testfile'
  to 'http://localhost:60000/httpd/d15900f5-34ee-4bba-bb10-73d60d1c0d75' with http...
Uploading 20 bytes... data sent, waiting... done.
- done in 0.0042 seconds.
'testfile' (20 bytes) uploaded as '/tmp/testfile'.
$ arc_storage_cli.py stat /tmp/testfile
- ARC_BARTENDER_URL environment variable not found, using http://localhost:60000/Bartender
- Calling the Bartender's stat method...
- done in 0.05 seconds.
'/tmp/testfile': found
states
  checksumType: md5
  neededReplicas: 1
  size: 20
  checksum: 9a9dffa22d227afe0f1959f936993a80
timestamps
  created: 1210232135.57
parents
  51e12fab-fd3d-43ec-9bc5-17041da3f0b2/testfile: parent
locations
  http://localhost:60000/Shepherd fc0d3d99-6406-4c43-b2eb-c7ec6d6ab7fe: alive
entry
  type: file
$ arc_storage_cli.py get /tmp/testfile newfile
- ARC_BARTENDER_URL environment variable not found, using http://localhost:60000/Bartender
- Calling the Bartender's getFile method...
- done in 0.05 seconds.
- Got transfer URL: http://localhost:60000/httpd/dab911d0-110f-468e-b0c3-627af6e3af31
- Downloading from 'http://localhost:60000/httpd/dab911d0-110f-468e-b0c3-627af6e3af31'
```

```

    to 'newfile' with http...
Downloading 20 bytes... done.
- done in 0.0035 seconds.
'/tmp/testfile' (20 bytes) downloaded as 'newfile'.
$ cat newfile
This is a testfile.

```

You can find more examples in the SVN<sup>3</sup>.

There are plans to create more sophisticated CLI and GUI tools and to create a FUSE<sup>4</sup> module and Windows Shell Extensions<sup>5</sup> to be able to mount the ARC1 storage namespace into the local filesystem namespace, and use it with the commands of the operating system.

## 3.6 Other plans

### 3.6.1 Integrating third-party storage solutions

To integrate existing files on a third-party storage to our namespace thus make them accessible through the interface of the ARC1 storage, we need a service which provides a common interface to the Bartender services, and hides the details of accessing the different third-party storages. It should translate the method calls, the *gets*, *puts* and *removes* and the *ACL*<sup>6</sup> modifications, and try to create transform metadata according to the data model of the ARC1 storage (Section 3.2.3). The files here would be referenced by a path which is local in the namespace of the third-party storage. The interface could be something like this:

**get(getRequest)** returns list of (*requestID*, *getResponseData*) very similar to the *get* method of the Shepherd

**put(putRequest)** returns a list of (*requestID*, *putResponseData*) very similar to the *put* method of the Shepherd

**delete(deleteRequest)** returns a list of (*requestID*, *success*) very similar to the *delete* method of the Shepherd

**stat(statRequest)** returns a list of (*requestID*, *statResponse*) very similar to the *stat* method of the Shepherd

**list(listRequest)** returns *listResponse* very similar to the *list* method of the Bartender

**move(moveRequest)** returns a list of (*requestID*, *status*) very similar to the *move* method of the Bartender

This interface has some methods similar to the Shepherd and some other methods similar to the Bartender.

### 3.6.2 Security

These are the planned actions which can be used for access control:

- *read*: user can get the list of entries in the collection; user can download the file
- *addEntry*: user can add a new entry to the collection;
- *removeEntry*: user can remove any entry from the collection
- *delete*: user can delete the collection if it is empty; user can delete a file (if you want to remove a file/collection, then the Bartender needs to remove the entry from the parent collection, and then delete the file/collection itself, so you need to have both permissions)

<sup>3</sup><http://svn.nordugrid.org/trac/nordugrid/browser/arc1/trunk/src/services/storage/README>

<sup>4</sup>Filesystem in Userspace, <http://fuse.sourceforge.net/>

<sup>5</sup><http://msdn.microsoft.com/en-us/magazine/cc188741.aspx>

<sup>6</sup>Access Control Lists, policies



- *modifyPolicy*: user can modify the policy of the file/collection
- *modifyStates*: user can modify some special metadata of the file/collection (close the collection, change the number of needed replica of the file)
- *modifyMetadata*: user can modify the arbitrary metadata section of the file/collection (these are key-value pairs)

When a user has the permission in the Librarian to download a file then the user should have permission to access at least one of the file's replica, so there should be a Shepherd which allows the user to get the file. If the Bartender has permission to access the Shepherd, then the Bartender should create an assertion which allows the user to access the file. This could be a signed token which contains a policy defining access to particular file. But all these are currently just plans.